

# Machine Learning

## Clustering - Beyond KMeans

Claudio Sartori

DISI

Department of Computer Science and Engineering – University of Bologna, Italy

[claudio.sartori@unibo.it](mailto:claudio.sartori@unibo.it)

|   |                                          |    |
|---|------------------------------------------|----|
| 1 | Hierarchical clustering                  | 2  |
|   | ● Separation                             | 5  |
|   | ● Single linkage hierarchical clustering | 7  |
|   | ● Examples                               | 15 |
| 2 | Density based clustering                 | 27 |
| 3 | Model based clustering                   | 47 |
| 4 | Final remarks                            | 65 |

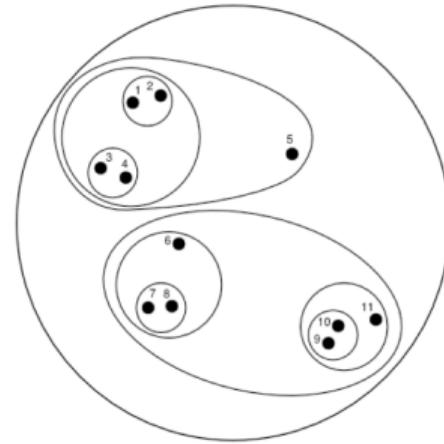
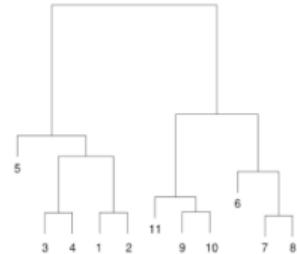
# Hierarchical clustering

Generates a **nested structure** of clusters

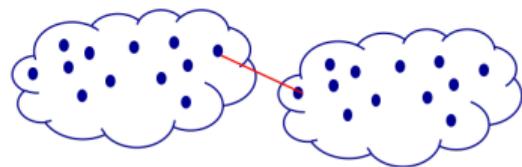
- Agglomerative (bottom up)
  - as a starting state, each data point is a cluster
  - in each step the two **less separated** clusters are merged into one
  - a measure of **separation between clusters** is needed
- Divisive (top down)
  - as a starting state, the entire dataset is the only cluster
  - in each step, the cluster with the lowest cohesion is split
  - a measure of cluster cohesion and a split procedure are needed

# Hierarchical clustering output

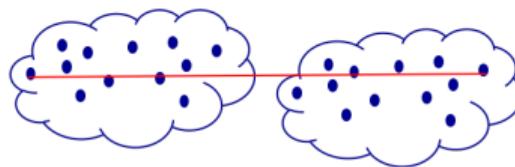
- Dendrogram (left)
- Nested cluster diagram (right)
- They represent the same structure
- The representation is the same for agglomerative and divisive
- The agglomerative methods are the most used



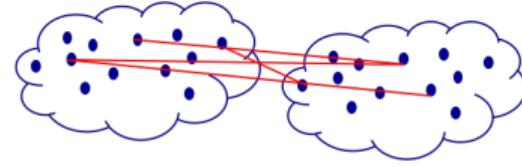
# Separation between clusters – Graph based



Single Link



Complete Link



Average Link

$$\text{Sep}(k_i, k_j) = \min_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

$$\text{Sep}(k_i, k_j) = \max_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

$$\text{Sep}(k_i, k_j) = \frac{1}{|k_i||k_j|} \sum_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

The distance between sets is based on the distances between objects belonging to the two sets, respectively

# Separation between clusters – Ward's method

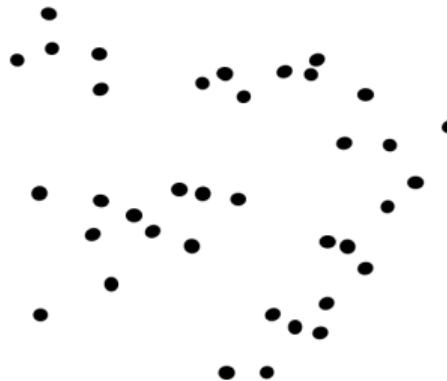
Given two sets  $S_1$  and  $S_2$  with the respective  $SSE(S_1)$  and  $SSE(S_2)$ , the distance is computed as

$$d(S_1, S_2) = SSE(S_1 \cup S_2) - (SSE(S_1) + SSE(S_2))$$

Smaller separation implies a lower increase in the SSE after merging

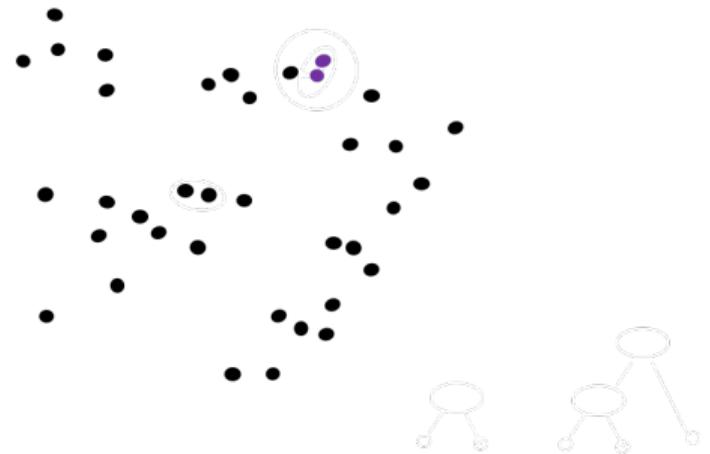
# Single linkage hierarchical clustering I

1. Initialization: every object is a cluster



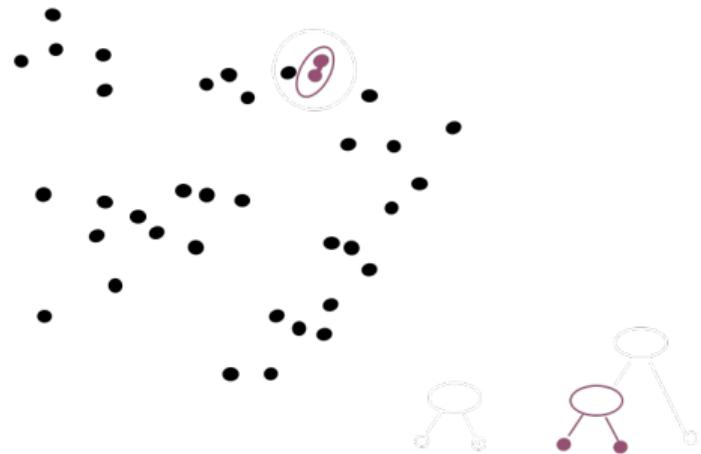
# Single linkage hierarchical clustering II

1. Initialization: every object is a cluster
2. Find the **less separated** pair



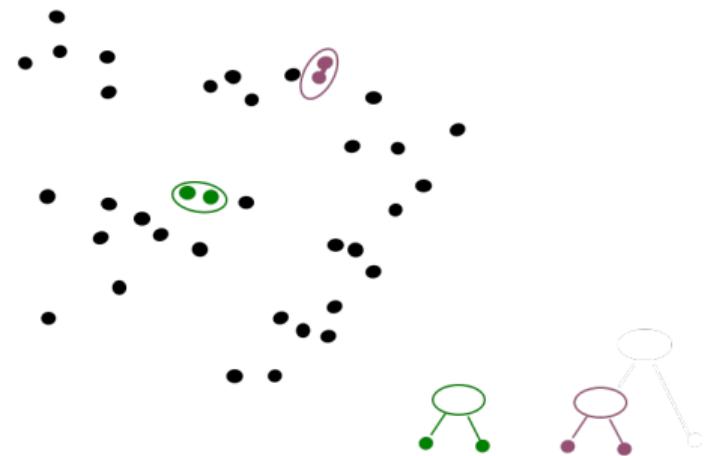
# Single linkage hierarchical clustering III

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster



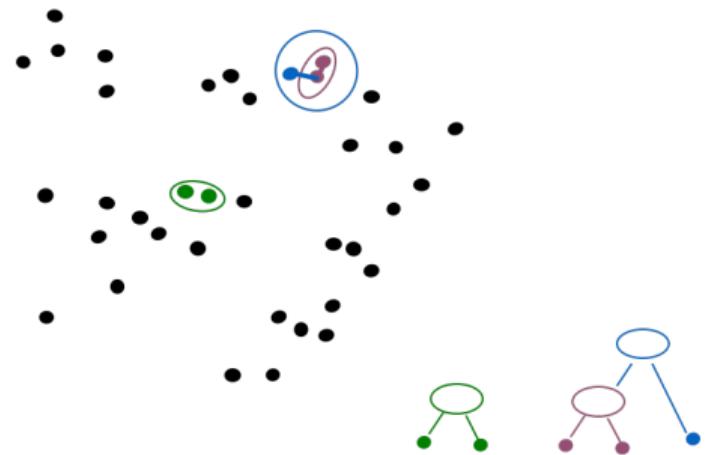
# Single linkage hierarchical clustering IV

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat



# Single linkage hierarchical clustering V

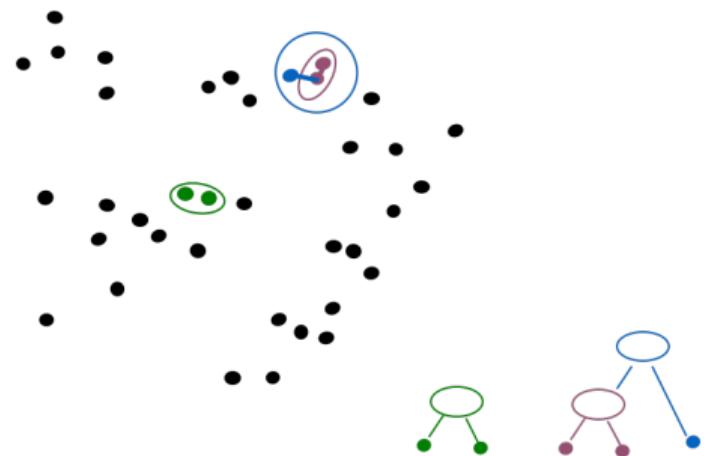
1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat ...



# Single linkage hierarchical clustering VI

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat ...

The result is a **dendrogram** (taxonomy, object hierarchy)



# Single linkage algorithm

- Initialize the clusters, one for each objects
- Compute the **distance matrix** between the clusters, squared, symmetric, the size is the number of objects  $N$ , the main diagonal is null
- While the number of clusters is greater than 1
  - find the two clusters with lowest separation, say  $k_r$  and  $k_s$
  - merge them in a cluster
  - delete from the distance matrix the rows and columns  $r$  and  $s$  and insert one new row and column with the distances of the new cluster from the others

$$\text{Dist}(k_k, k_{(r+s)}) = \min(\text{Dist}(k_k, k_r), \text{Dist}(k_k, k_s)) \forall k \in [1, K]$$

# Time and space complexity

- Space and time:  $\mathcal{O}(N^2)$  for the computation and the storage of the distance matrix
- Worst case  $N - 1$  iterations to reach the final single cluster
- For the  $i$ -th step of the main iteration:
  - search of the pair to merge  $\mathcal{O}((N - i)^2)$
  - recomputation of the distance matrix  $\mathcal{O}((N - i))$
- Time, in summary:  $\mathcal{O}(N^3)$ 
  - can be reduced to  $\mathcal{O}(N^2 \log(N))$  with indexing structures

# Italian cities example I

|    | BA  | FI  | MI  | NA  | RM  | TO  |
|----|-----|-----|-----|-----|-----|-----|
| BA | 0   | 662 | 877 | 255 | 412 | 996 |
| FI | 662 | 0   | 295 | 468 | 268 | 400 |
| MI | 877 | 295 | 0   | 754 | 564 | 138 |
| NA | 255 | 468 | 754 | 0   | 219 | 869 |
| RM | 412 | 268 | 564 | 219 | 0   | 669 |
| TO | 996 | 400 | 138 | 869 | 669 | 0   |



# Italian cities example II

|    | BA  | FI  | MI  | NA  | RM  | TO  |
|----|-----|-----|-----|-----|-----|-----|
| BA | 0   | 662 | 877 | 255 | 412 | 996 |
| FI | 662 | 0   | 295 | 468 | 268 | 400 |
| MI | 877 | 295 | 0   | 754 | 564 | 138 |
| NA | 255 | 468 | 754 | 0   | 219 | 869 |
| RM | 412 | 268 | 564 | 219 | 0   | 669 |
| TO | 996 | 400 | 138 | 869 | 669 | 0   |



# Italian cities example III

|       | BA  | FI  | MI/TO | NA  | RM  |
|-------|-----|-----|-------|-----|-----|
| BA    | 0   | 662 | 877   | 255 | 412 |
| FI    | 662 | 0   | 295   | 468 | 268 |
| MI/TO | 877 | 295 | 0     | 754 | 564 |
| NA    | 255 | 468 | 754   | 0   | 219 |
| RM    | 412 | 268 | 564   | 219 | 0   |



# Italian cities example IV

|       | BA  | FI  | MI/TO | NA/RM |
|-------|-----|-----|-------|-------|
| BA    | 0   | 662 | 877   | 255   |
| FI    | 662 | 0   | 295   | 268   |
| MI/TO | 877 | 295 | 0     | 564   |
| NA/RM | 255 | 268 | 564   | 0     |



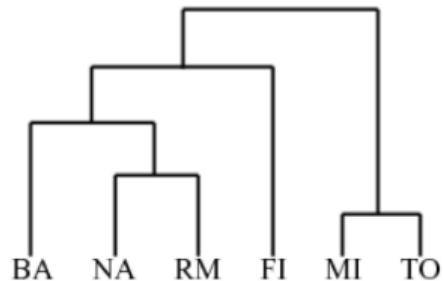
# Italian cities example V

|          | BA/NA/RM | FI  | MI/TO |
|----------|----------|-----|-------|
| BA/NA/RM | 0        | 268 | 564   |
| FI       | 268      | 0   | 295   |
| MI/TO    | 564      | 295 | 0     |

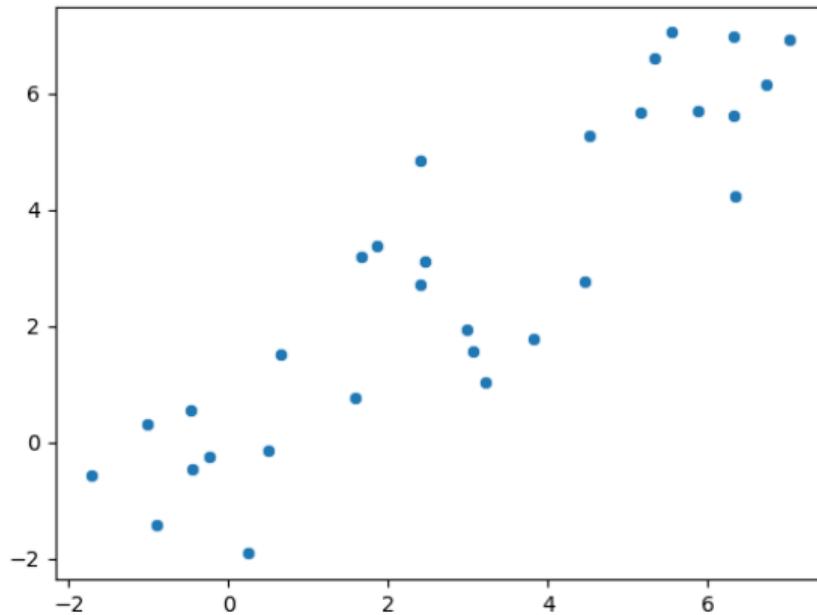


# Italian cities example VI

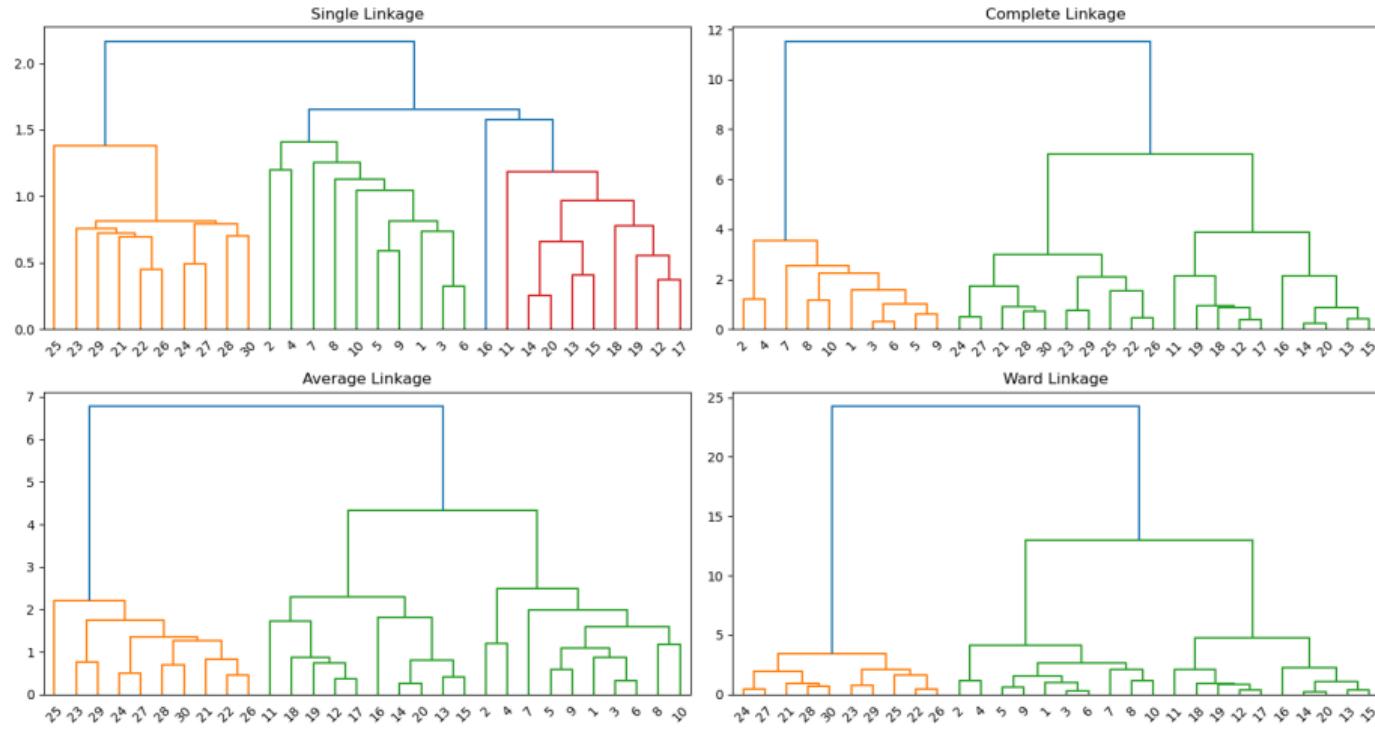
|             | BA/FI/NA/RM | MI/TO |
|-------------|-------------|-------|
| BA/FI/NA/RM | 0           | 295   |
| MI/TO       | 295         | 0     |



# Comparison of linking methods - Sample data

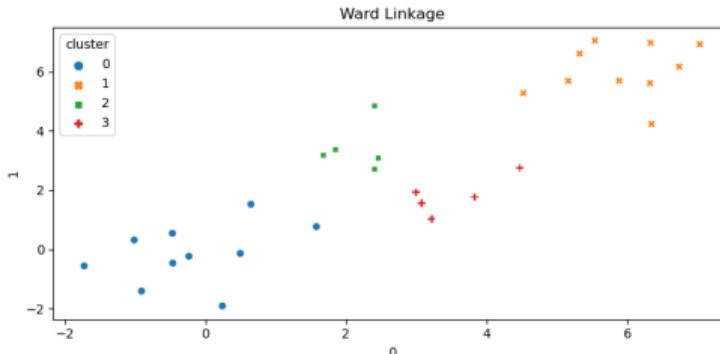
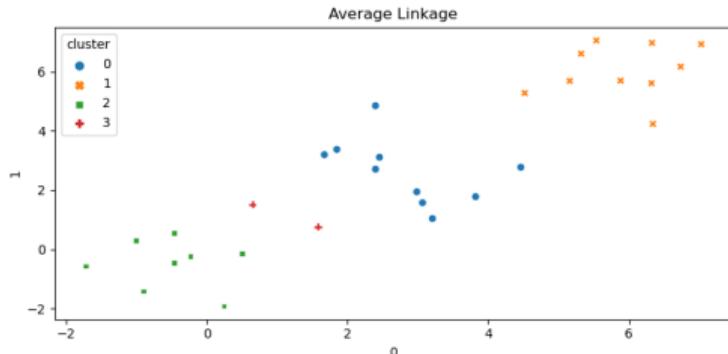
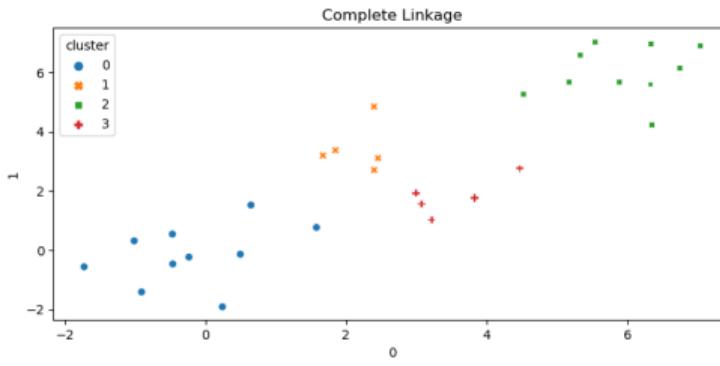
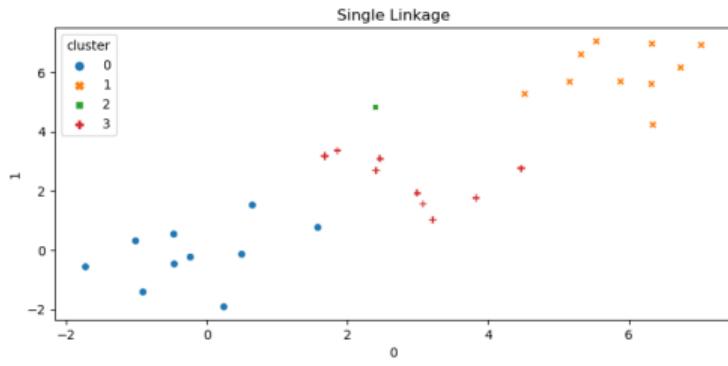


# Comparison of linking methods - Dendograms



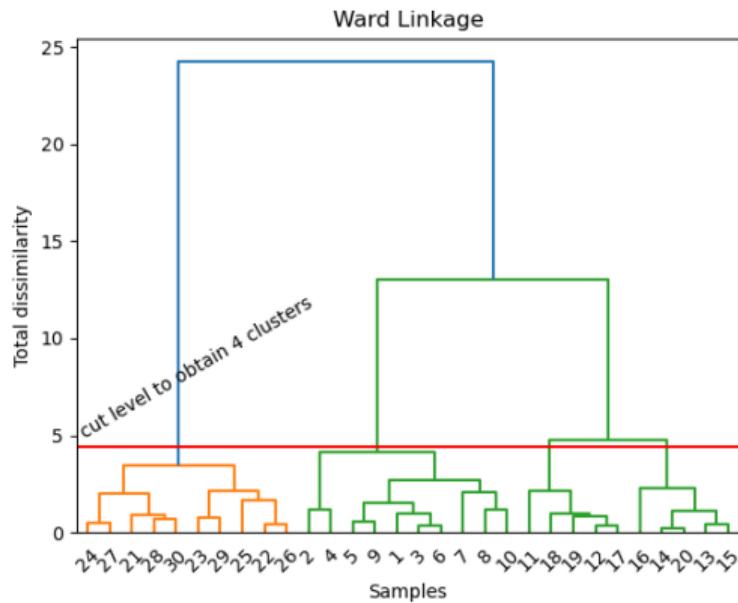
# Comparison of linking methods

Labelled data with 4 clusters



# Generating the clustering scheme

- The desired clustering scheme is obtained by choosing the number of clusters
- The choice of the level is application dependent, and can also be guided by indexes, as in the case of K-means
- The choice is equivalent to **cutting** the dendrogram at the appropriate level



# Discussion I

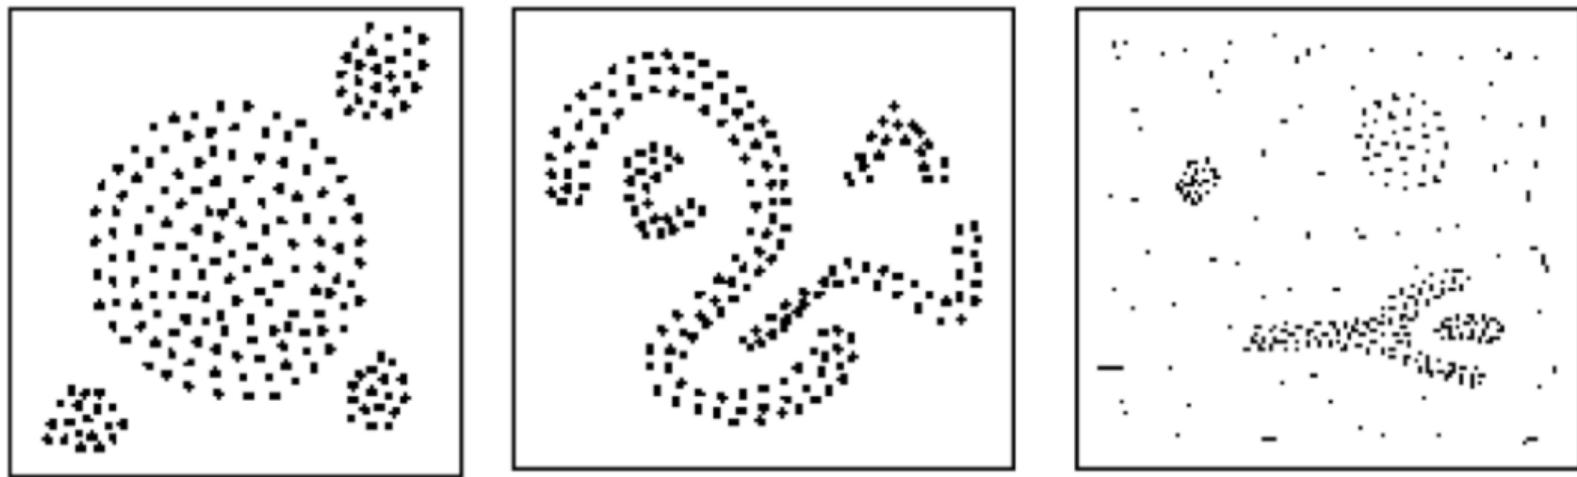
- The vertical axis in the dendrogram is the **total dissimilarity** inside the clusters, which obviously increases for decreasing number of clusters
- The **diameter** of a cluster is the distance among the most separated objects
  - Single linkage tends to generate clusters with larger diameters also at low levels
  - Complete linkage tends to generate more compact clusters
  - Ward linkage tends to generate more **spheric** clusters

# Discussion II

- :( The scaling is poor, due to the high complexity
- :( There isn't a global objective function, the decision is always local and cannot be undone
- : The dendrogram structure is of great help for the interpretation of the result
- : Empirically, the result is frequently good

|   |                          |    |
|---|--------------------------|----|
| 1 | Hierarchical clustering  | 2  |
| 2 | Density based clustering | 27 |
|   | ● DBSCAN                 | 30 |
| 3 | Model based clustering   | 47 |
| 4 | Final remarks            | 65 |

# Density based clustering



Clusters are high-density regions separated by low-density regions

# Computing density

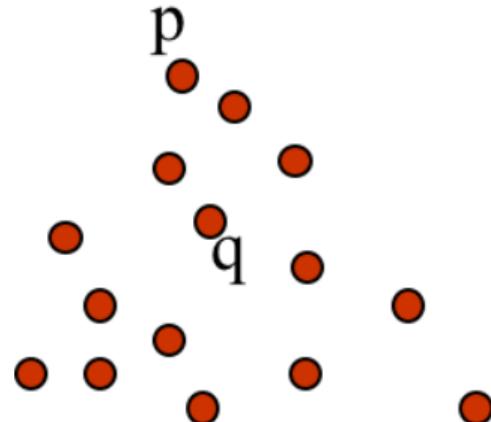
The two most obvious solutions

- Grid-based
  - split the (hyper)space into a regularly spaced grid
  - count the number of objects inside each grid element
- Object-centered
  - define the radius of a (hyper)sphere
  - attach to each object the number of objects which are inside that sphere

# DBSCAN – Density Based Spatial Clustering of Applications with Noise<sup>1</sup>

## Intuition

- intuitively,  $p$  is a border point, while  $q$  is a core point



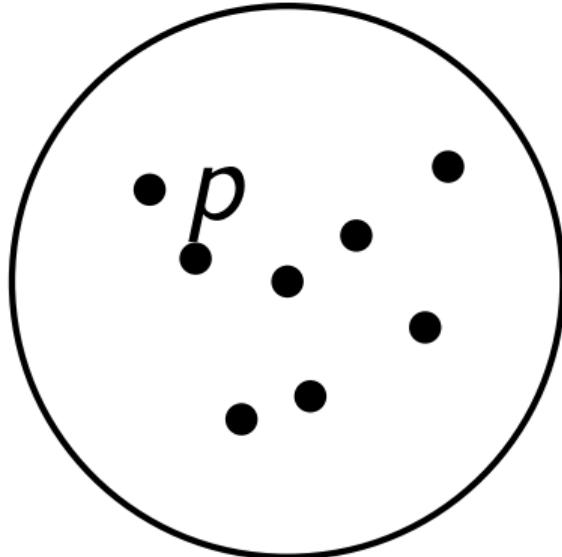
<sup>1</sup> [Ester et al.(1996)] Ester, Kriegel, Sander, and Xu]

# Formalisation of density vs sparsity

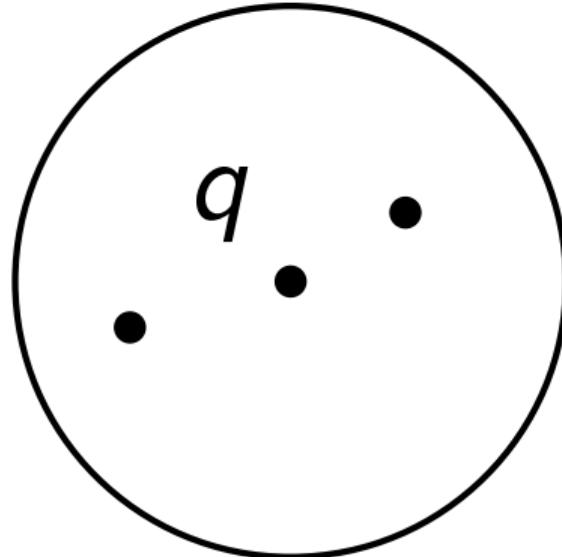
- define the neighbourhood of a point  $p$
- idea: define a hypersphere of radius  $\epsilon$  centered in the  $p$
- count how many neighbours are in the hypersphere  $|N_\epsilon(p)|$
- define a threshold

# Eps-Neighborhoods: Dense vs Sparse Regions

Dense Region

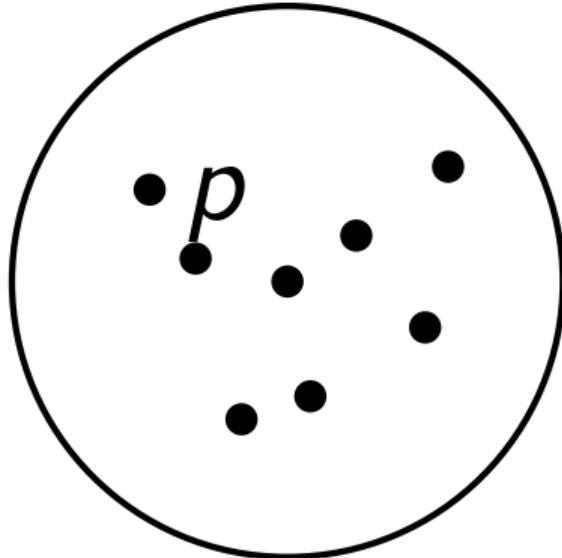


Sparse Region

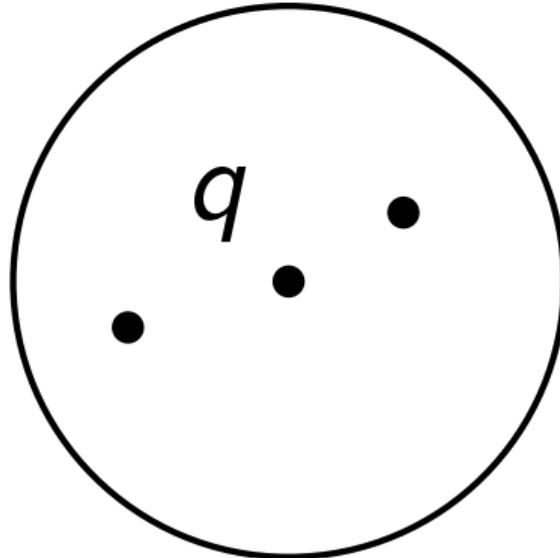


# Eps-Neighborhoods: Dense vs Sparse Regions

Dense Region



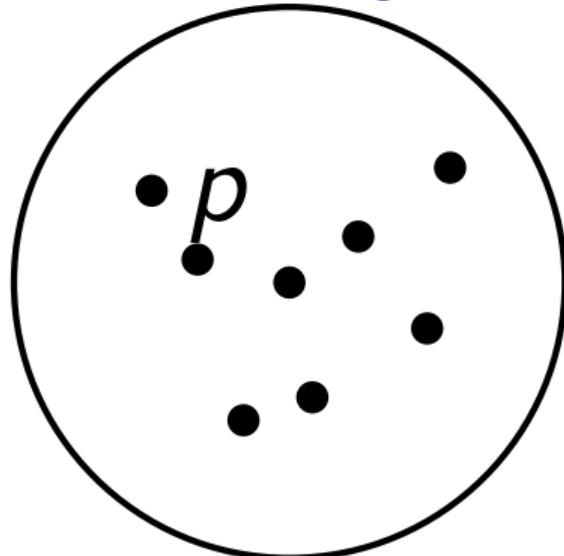
Sparse Region



$|N_\epsilon(p)| \geq \text{MinPts} \Rightarrow p$  is a core point.

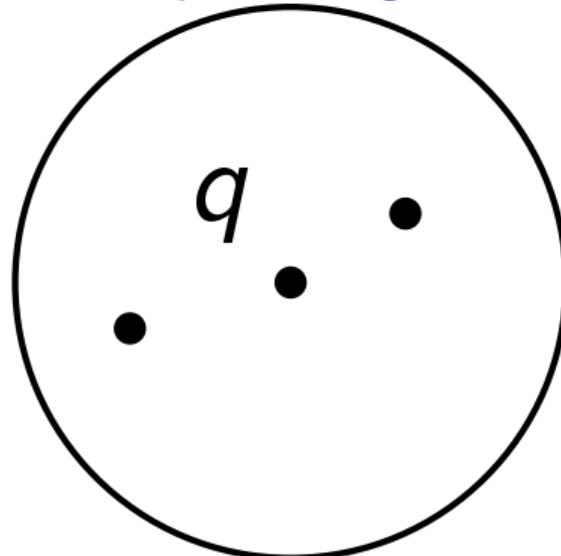
# Eps-Neighborhoods: Dense vs Sparse Regions

Dense Region



$|N_\varepsilon(p)| \geq \text{MinPts} \Rightarrow p$  is a **core point**.

Sparse Region

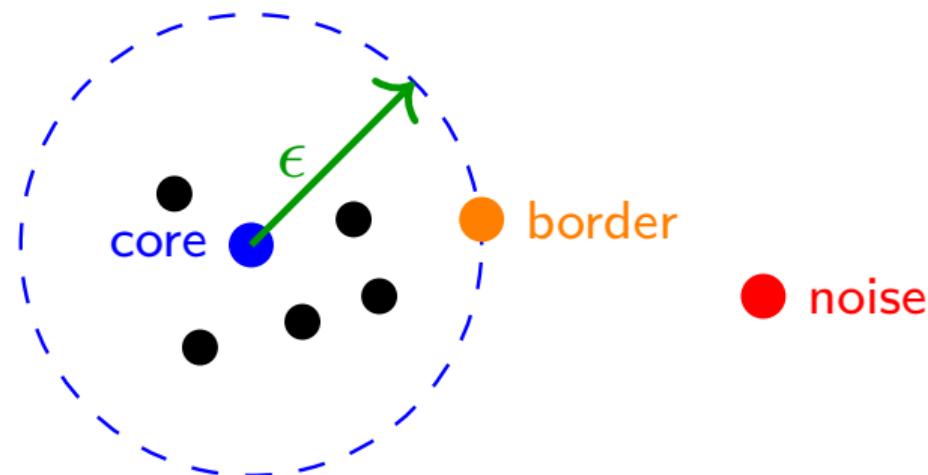


$|N_\varepsilon(q)| < \text{MinPts} \Rightarrow q$  is **not** a core point.

# Core, Border, and Noise Points

## Definitions

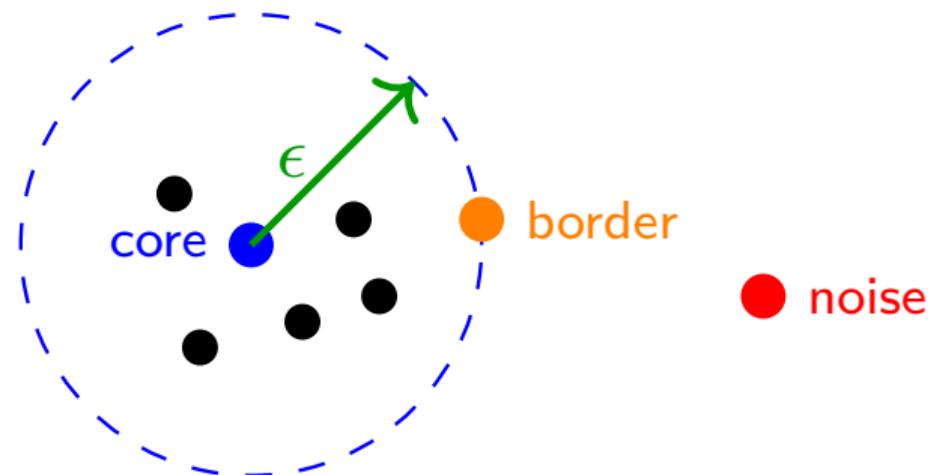
- **Core point:** at least  $\text{MinPts}$  points within  $\varepsilon$ .



# Core, Border, and Noise Points

## Definitions

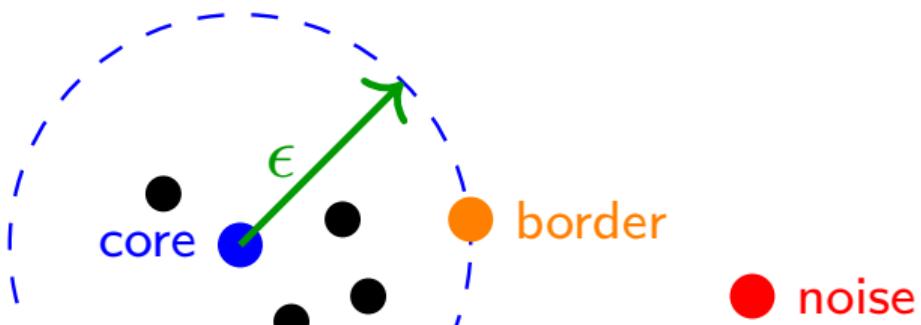
- **Core point:** at least  $\text{MinPts}$  points within  $\varepsilon$ .
- **Border point:** not a core point, but neighbor of a core.



# Core, Border, and Noise Points

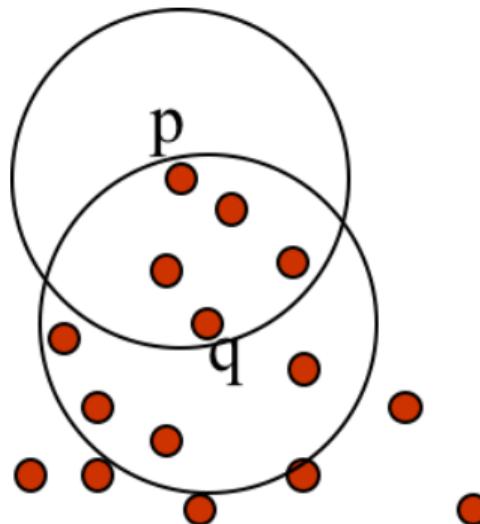
## Definitions

- **Core point:** at least  $\text{MinPts}$  points within  $\varepsilon$ .
- **Border point:** not a core point, but neighbor of a core.
- **Noise point:** neither core nor border.



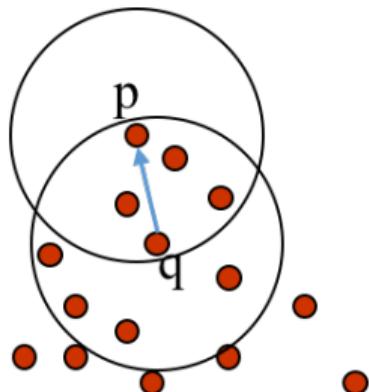
# Neighborhood

- define a radius  $\epsilon$  and define as **neighborhood** of a point the  $\epsilon$ -hypersphere centered at that point
- points  $p$  and  $q$  are one in the neighborhood of the other
  - neighborhood is **symmetric**



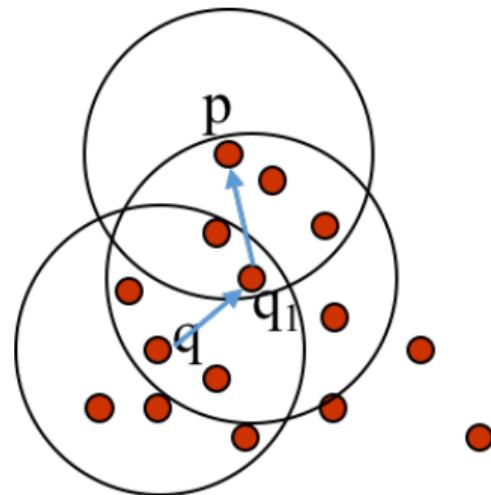
# Direct Density Reachability

- define a threshold `minPoints` and define as **core** a point with at least `minPoints` points in its neighborhood, as **border** otherwise
  - with  $\text{minPoints} = 5$ ,  $q$  is core,  $p$  is border
- define that a point  $p$  is **directly density reachable** from point  $q$  iff
  - $q$  is core
  - $q$  is in the neighborhood of  $p$
- direct density reachability is not symmetric
  - in the example  $q$  is not directly density reachable from  $p$ , since  $p$  is border



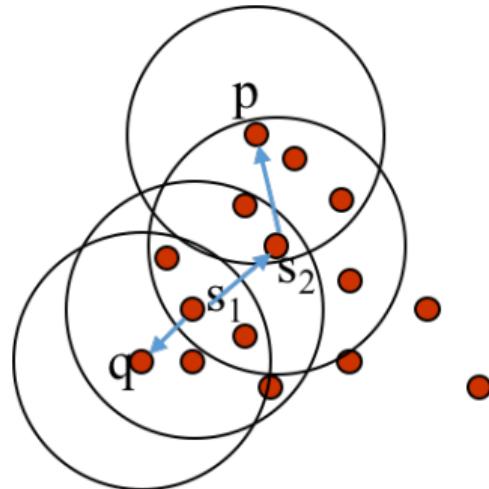
# Density Reachability

- a point  $p$  is **density reachable** from point  $q$  iff
  - $q$  is core
  - there is a sequence of points point  $q_i$  such that  $q_{i+1}$  is directly density reachable from  $q_i$ ,  $i \in [1, nq]$ ,  $q_1$  is directly reachable from  $q$  and  $p$  is directly density reachable from  $q_{nq}$
- reachability is not symmetric
  - in the example  $q$  is not density reachable from  $p$ , since  $p$  is border



# Density Connection

- a point  $p$  is **density connected** to point  $q$  iff there is a point  $s$  such that  $p$  and  $q$  are density reachable from  $s$
- density connection is symmetric



# Generation of clusters

- A **cluster** is a maximal set of points connected by **density**
- Border points which are not connected by density to any core point are labelled as **noise**

# DBSCAN Main Loop

**Input:** SetOfPoints: UNCLASSIFIED  
points; Eps; MinPts

**Output:** SetOfPoints

```
ClusterId ← nextId(NOISE);  
for  $i \leftarrow 1$  to SetOfPoints.size do  
    Point ← SetOfPoints.get(i);  
    if Point.CId = UNCLASSIFIED then  
        if ExpandCluster(SetOfPoints,  
                           Point, ClusterId, Eps, MinPts)  
        then  
            ClusterId ←  
            nextId(ClusterId);
```

## Explanation:

- All points start as UNCLASSIFIED.

# DBSCAN Main Loop

**Input:** SetOfPoints: UNCLASSIFIED  
points; Eps; MinPts

**Output:** SetOfPoints

```
ClusterId ← nextId(NOISE);  
for  $i \leftarrow 1$  to SetOfPoints.size do  
    Point ← SetOfPoints.get(i);  
    if Point.CId = UNCLASSIFIED then  
        if ExpandCluster(SetOfPoints,  
                           Point, ClusterId, Eps, MinPts)  
        then  
            ClusterId ←  
            nextId(ClusterId);
```

## Explanation:

- All points start as UNCLASSIFIED.
- We iterate through the dataset exactly once.

# DBSCAN Main Loop

**Input:** SetOfPoints: UNCLASSIFIED  
points; Eps; MinPts

**Output:** SetOfPoints

```
ClusterId ← nextId(NOISE);
for i ← 1 to SetOfPoints.size do
    Point ← SetOfPoints.get(i);
    if Point.CId = UNCLASSIFIED then
        if ExpandCluster(SetOfPoints,
                        Point, ClusterId, Eps, MinPts)
        then
            ClusterId ←
                nextId(ClusterId);
```

## Explanation:

- All points start as UNCLASSIFIED.
- We iterate through the dataset exactly once.
- Each unclassified point attempts to start a cluster.

# DBSCAN Main Loop

**Input:** SetOfPoints: UNCLASSIFIED  
points; Eps; MinPts

**Output:** SetOfPoints

```
ClusterId ← nextId(NOISE);
for i ← 1 to SetOfPoints.size do
    Point ← SetOfPoints.get(i);
    if Point.CId = UNCLASSIFIED then
        if ExpandCluster(SetOfPoints,
                        Point, ClusterId, Eps, MinPts)
        then
            ClusterId ←
                nextId(ClusterId);
```

## Explanation:

- All points start as UNCLASSIFIED.
- We iterate through the dataset exactly once.
- Each unclassified point attempts to start a cluster.
- If successful, the cluster grows via ExpandCluster.

# DBSCAN Main Loop

**Input:** SetOfPoints: UNCLASSIFIED  
points; Eps; MinPts

**Output:** SetOfPoints

```
ClusterId ← nextId(NOISE);
for i ← 1 to SetOfPoints.size do
    Point ← SetOfPoints.get(i);
    if Point.CId = UNCLASSIFIED then
        if ExpandCluster(SetOfPoints,
                        Point, ClusterId, Eps, MinPts)
        then
            ClusterId ←
                nextId(ClusterId);
```

## Explanation:

- All points start as UNCLASSIFIED.
- We iterate through the dataset exactly once.
- Each unclassified point attempts to start a cluster.
- If successful, the cluster grows via ExpandCluster.
- ClusterId increments only after a successful expansion.

# ExpandCluster Subroutine

---

**Input:** SetOfPoints, Point, ClusterId, Eps, MinPts

**Output:** True if cluster expanded, False otherwise

$\text{NeighborPts} \leftarrow \text{RegionQuery}(\text{SetOfPoints}, \text{Point}, \text{Eps});$

# ExpandCluster Subroutine

---

**Input:** SetOfPoints, Point, ClusterId, Eps, MinPts

**Output:** True if cluster expanded, False otherwise

*NeighborPts*  $\leftarrow$  RegionQuery(SetOfPoints, Point, Eps);

**if**  $|NeighborPts| < MinPts$  **then**

*Point.ClId*  $\leftarrow$  NOISE;

**return** False;

# ExpandCluster Subroutine

---

**Input:** SetOfPoints, Point, ClusterId, Eps, MinPts

**Output:** True if cluster expanded, False otherwise

*NeighborPts*  $\leftarrow$  RegionQuery(SetOfPoints, Point, Eps);

**if**  $|NeighborPts| < MinPts$  **then**

*Point.CId*  $\leftarrow$  NOISE;

**return** False;

Assign *ClusterId* to all points in *NeighborPts*;

# ExpandCluster Subroutine

---

**Input:** SetOfPoints, Point, ClusterId, Eps, MinPts

**Output:** True if cluster expanded, False otherwise

$\text{NeighborPts} \leftarrow \text{RegionQuery}(\text{SetOfPoints}, \text{Point}, \text{Eps});$

**if**  $|\text{NeighborPts}| < \text{MinPts}$  **then**

$\text{Point.ClId} \leftarrow \text{NOISE};$

**return** False;

Assign  $\text{ClusterId}$  to all points in  $\text{NeighborPts}$ ;

**for** each  $P'$  in  $\text{NeighborPts}$  **do**

$\text{NeighborPts}' \leftarrow \text{RegionQuery}(\text{SetOfPoints}, P', \text{Eps});$

**if**  $|\text{NeighborPts}'| \geq \text{MinPts}$  **then**

        merge( $\text{NeighborPts}$ ,  $\text{NeighborPts}'$ );

# ExpandCluster Subroutine

---

**Input:** SetOfPoints, Point, ClusterId, Eps, MinPts

**Output:** True if cluster expanded, False otherwise

$\text{NeighborPts} \leftarrow \text{RegionQuery}(\text{SetOfPoints}, \text{Point}, \text{Eps});$

**if**  $|\text{NeighborPts}| < \text{MinPts}$  **then**

$\text{Point.ClId} \leftarrow \text{NOISE};$

**return** False;

Assign  $\text{ClusterId}$  to all points in  $\text{NeighborPts}$ ;

**for** each  $P'$  in  $\text{NeighborPts}$  **do**

$\text{NeighborPts}' \leftarrow \text{RegionQuery}(\text{SetOfPoints}, P', \text{Eps});$

**if**  $|\text{NeighborPts}'| \geq \text{MinPts}$  **then**

        merge( $\text{NeighborPts}$ ,  $\text{NeighborPts}'$ );

**return** True;

# RegionQuery Subroutine

---

---

**Input:** SetOfPoints, Point, Eps

**Output:** All points within distance  $\leq Eps$  of Point

$Neighbors \leftarrow \emptyset;$

---

# RegionQuery Subroutine

---

---

**Input:** SetOfPoints, Point, Eps

**Output:** All points within distance  $\leq Eps$  of Point

$Neighbors \leftarrow \emptyset;$

**for** each  $P'$  in SetOfPoints **do**

**if**  $distance(Point, P') \leq Eps$  **then**

        append  $P'$  to  $Neighbors$ ;

# RegionQuery Subroutine

---

**Input:** SetOfPoints, Point, Eps

**Output:** All points within distance  $\leq Eps$  of Point

*Neighbors*  $\leftarrow \emptyset$ ;

**for** each  $P'$  in SetOfPoints **do**

**if**  $\text{distance}(Point, P') \leq Eps$  **then**

append  $P'$  to *Neighbors*;

---

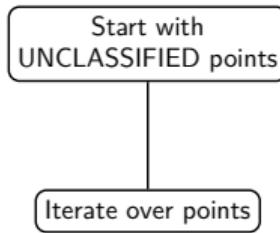
**return** *Neighbors*;

# DBSCAN: Full Algorithm Flow

Start with  
UNCLASSIFIED points

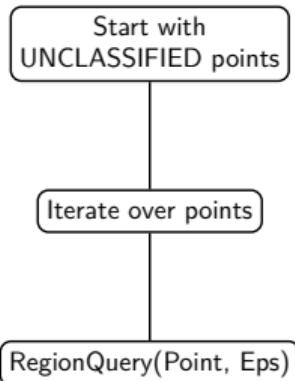
- Initialize all points as UNCLASSIFIED.

# DBSCAN: Full Algorithm Flow



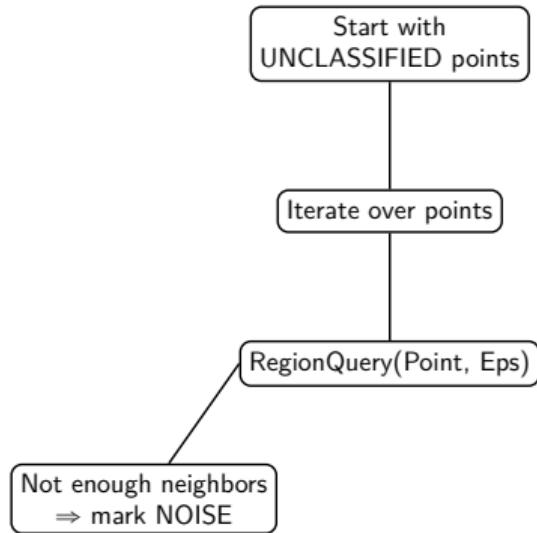
- Initialize all points as UNCLASSIFIED.
- Loop through the dataset.

# DBSCAN: Full Algorithm Flow



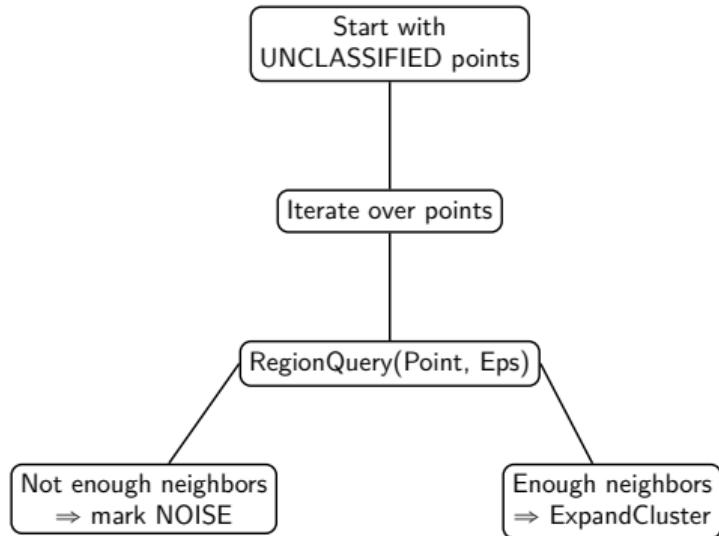
- Initialize all points as UNCLASSIFIED.
- Loop through the dataset.
- Perform RegionQuery for each unclassified point.

# DBSCAN: Full Algorithm Flow



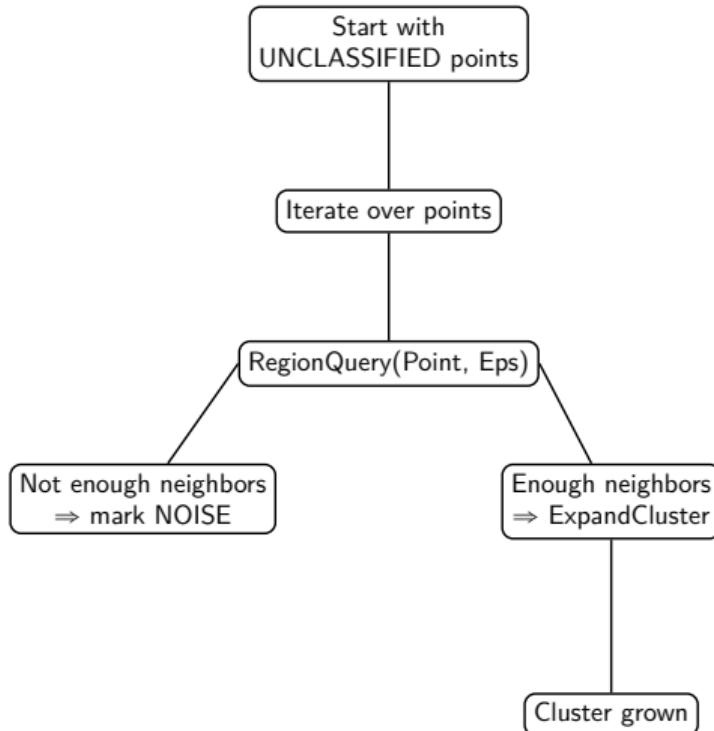
- Initialize all points as UNCLASSIFIED.
- Loop through the dataset.
- Perform RegionQuery for each unclassified point.
- If too few neighbors ? point becomes NOISE (or border).

# DBSCAN: Full Algorithm Flow



- Initialize all points as UNCLASSIFIED.
- Loop through the dataset.
- Perform RegionQuery for each unclassified point.
- If too few neighbors ? point becomes NOISE (or border).
- Otherwise start a cluster and expand it.

# DBSCAN: Full Algorithm Flow



- Initialize all points as UNCLASSIFIED.
- Loop through the dataset.
- Perform RegionQuery for each unclassified point.
- If too few neighbors ? point becomes NOISE (or border).
- Otherwise start a cluster and expand it.
- Cluster grows by recursively merging dense neighborhoods.

# How to set $\epsilon$ and minPoints?

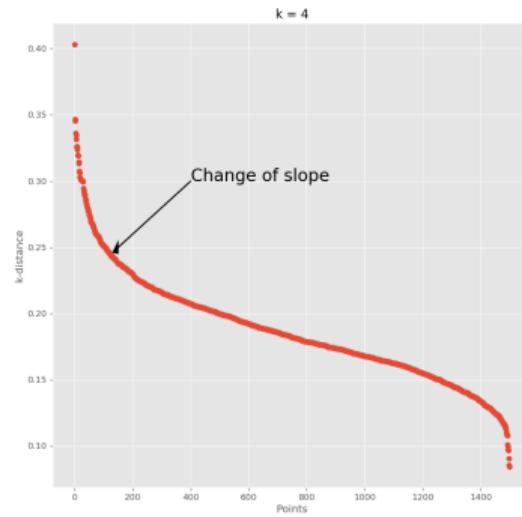
- As in many other machine learning algorithms, a **grid search** over several combination of hyperparameters can be useful
- As a **rule of thumb**, you can try  $\text{minPoints} = 2 * D$ , the number of dimensions
- Noise suggest an increase in  $\text{minPoints}$
- A guess for  $\epsilon$  requires more effort, considering the distance of the  $k$ -nearest neighbour, with  $k = \text{minPoints}$

# Good guess for $\epsilon$ !

- Consider the vector of the  $k$ -distances
  - choose  $k$
  - for each point we compute the distance of its  $k$ -nearest neighbour and we sort the points for decreasing  $k$ -distance
- Choosing a given  $k$ -distance as  $\epsilon$ , it turns out that all the points with a  $k$ -distance bigger than  $\epsilon$  will be considered as **border**
  - in the figure of next page they are the points to the left of the vertical of the chosen  $\epsilon$

# Good guess for $\epsilon \parallel$

- Usually, datasets which exhibit some tendency to clustering exhibit also a **change of slope**
- The best  $\epsilon$  can be found with a grid search in the **area of the change of slope**
  - the figure refers to a dataset with 1500 points and with  $\text{minPoints}=4$
  - this figure suggests a fine tuning of  $\epsilon$  in the interval 0.2–0.3



# Comments

- 😊 Finds clusters of any shape
- 😊 Is robust w.r.t. noise
- 😢 Problems if clusters have widely varying densities
- Being based on distances between points, the complexity is  $\mathcal{O}(N^2)$ 
  - reduced to  $\mathcal{O}(N \log(N))$  if spatial indexes, such as R\*, are available
- Very sensitive to the values of  $\epsilon$  and minPoints
- Decreasing  $\epsilon$  and increasing minPoints reduces the cluster size and increases the number of noise points



|   |                               |    |
|---|-------------------------------|----|
| 1 | Hierarchical clustering       | 2  |
| 2 | Density based clustering      | 27 |
| 3 | <b>Model based clustering</b> | 47 |
|   | ● Gaussian Mixture            | 49 |
| 4 | Final remarks                 | 65 |

# Model based (or statistic based) clustering<sup>2</sup>

- Estimate the parameters of a statistical model to maximize the ability of the model to **explain the data**
- The main technique is to use the **mixture models**
  - view the data as a set of observation from a mixture of different probability distributions
- Usually, the base model is a multivariate normal
  - well-known, easy to work with, good results
- The estimation is usually done using the **maximum likelihood**
  - given a set of data  $\mathcal{X}$ , the probability of the data, regarded as a function of the parameters, is called a **likelihood function**
- Attributes are assumed to be random independent variables

2 [Tan et al.(2006) Tan, Steinbach, and Kumar], Section 9.2.2

# Gaussian Mixture

## a.k.a. Expectation Maximization – EM

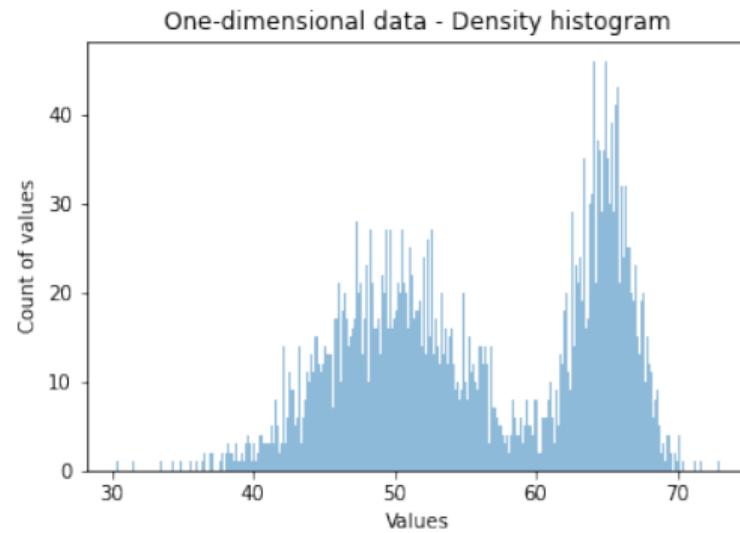
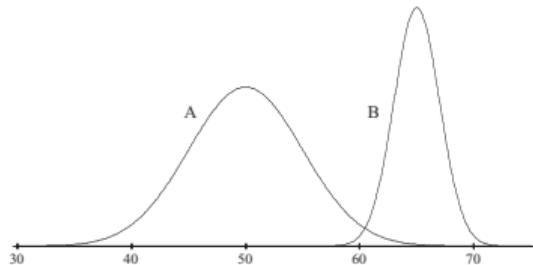
- If the data can be approximated by a single distribution, the derivation of the parameters is straightforward
- In the general case, with many mixed distributions, the EM algorithm is used

# EM algorithm

1. Select an initial set of model parameters
2. **repeat**
  - 2.1 **Expectation Step** – For each object, calculate the probability that each object belongs to each distribution
  - 2.2 **Maximization Step** – Given the probabilities from the expectation step, find the new estimates of the parameters that maximize the expected likelihood
3. **until** – the parameters do not change (or the change is below a specified threshold)

# One dimension mixture example

- Case with one dimension, two components
- Synthetic data randomly generated with two gaussians  
 $\mu_A = 50, \sigma_A = 5, p_A = 0.6$   
 $\mu_B = 65, \sigma_B = 2, p_B = 0.4$



# EM – one dimension, two clusters example I

- Need to estimate 5 parameters
  - mean and standard deviation for cluster A
  - mean and standard deviation for cluster B
  - sampling probability  $p$  for cluster A

$$\Pr(A|x) = \frac{\Pr(x|A)\Pr(A)}{\Pr(x)} = \frac{f(x; \mu_A, \sigma_A)p_A}{\Pr(x)}$$
$$f(x; \mu_A, \sigma_A) = \frac{1}{\sqrt{2\pi}\sigma} e^{\frac{(x-\mu)^2}{2\sigma^2}}$$

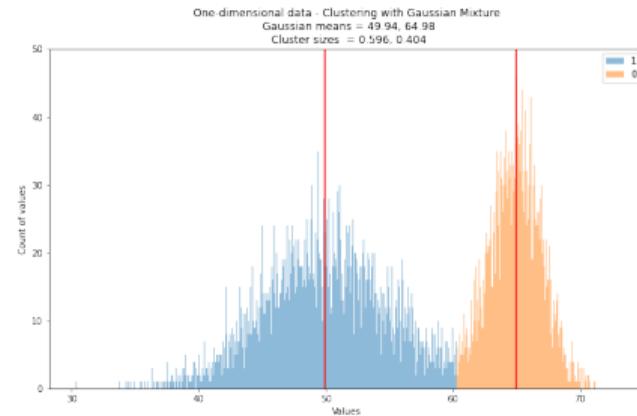
# EM – one dimension, two clusters example II

- Repeat until convergence
  - Expectation: Compute  $p_A$  and  $p_B$  using the current distribution parameters
    - Compute the numerators for  $\Pr(A|x)$  and  $\Pr(B|x)$  and normalize dividing by their sum
  - Maximization of the distribution likelihood given the data
    - Compute the new distributions parameters, weighting the probabilities according to the current distribution parameters
- After convergence label each object with  $A$  or  $B$  according to the maximum probability, given the last distribution parameters

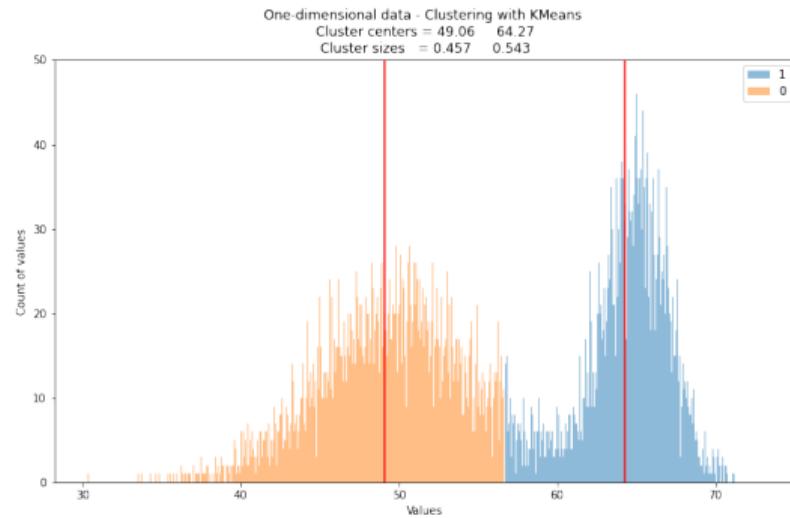
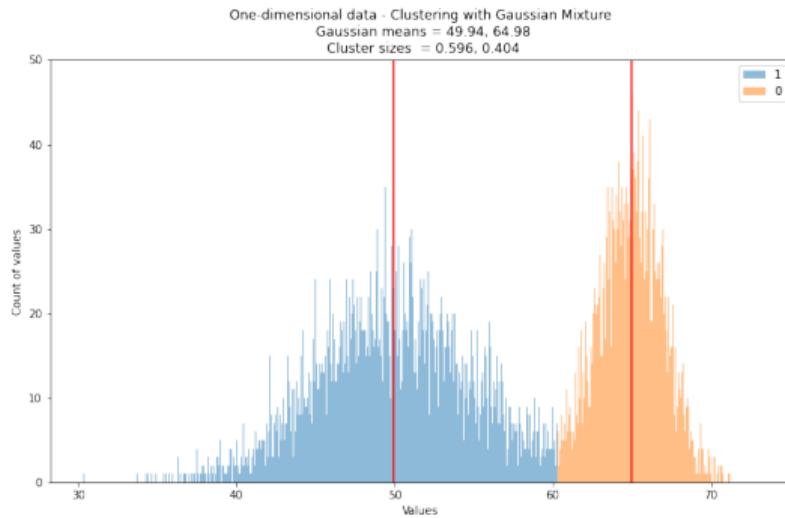
# One dimension example - Gaussian Mixture result

Estimated parameters of the distributions

|   | weight | mean      | deviation |
|---|--------|-----------|-----------|
| 0 | 0.4037 | 64.978754 | 2.030968  |
| 1 | 0.5963 | 49.939990 | 4.999235  |



# Gaussian Mixture and KMeans – Comparison of results



- These data have a **bimodal, gaussian-like** distribution
- The EM algorithm is founded on the hypothesis of modelling data with gaussians
- KMeans is **non-parametric**, and in this case the performance is **worse**

|   |                          |    |
|---|--------------------------|----|
| 1 | Hierarchical clustering  | 2  |
| 2 | Density based clustering | 27 |
| 3 | Model based clustering   | 47 |
| 4 | Final remarks            | 65 |

# BIRCH Clustering: Overview

**BIRCH** ([Balanced Iterative Reducing and Clustering using Hierarchies](#)) is a scalable clustering algorithm designed for [large datasets](#). It incrementally builds a compact representation of the data called a **CF-tree** ([Clustering Feature Tree](#)), which summarizes data points using:

- [N](#): number of points
- [LS](#): linear sum of points
- [SS](#): squared sum of points

A CF-tree enables fast, memory-efficient clustering with a single scan of the data.

# Key Idea: The CF Node

Each CF-tree node stores a set of *Clustering Features* that summarize a subcluster.

- A **CF** triple  $\langle N, LS, SS \rangle$  compactly captures centroid and radius.
- Internal nodes group subclusters hierarchically.
- Leaf nodes contain the final *subcluster summaries*.

This structure allows BIRCH to operate with constrained memory while preserving cluster quality.

# Algorithm Steps

The BIRCH algorithm proceeds in four conceptual phases:

1. **Build CF-tree:** Insert each point; absorb it into the closest subcluster if within **threshold  $T$** ; otherwise split.
2. **Condense CF-tree:** Optionally remove outliers or merge small subclusters.
3. **Global Clustering:** Apply a standard clustering method (e.g. **agglomerative**) to the leaf subclusters.
4. **Refinement:** Optionally reassign original data to improved cluster centers.

Steps 3 and 4 are optional but improve accuracy.

# Advantages of BIRCH

- Designed for **very large datasets** (single scan).
- Memory usage remains low via compact CF representations.
- Naturally supports incremental and dynamic updates.
- Performs well for **numerical, metric-space** data.

**Limitation:** BIRCH struggles with non-spherical or poorly separated clusters due to its reliance on centroid-based thresholding.

# Spectral Clustering: Overview

**Spectral clustering** uses the eigenstructure of a similarity graph to partition data into clusters.

It transforms the data into a low-dimensional space using the **graph Laplacian**, where clusters become more easily separable.

Works especially well for **non-convex** or **manifold-shaped** clusters.

# Similarity Graph

Given data points, construct a weighted graph  $G$ :

- Vertices = data points
- Edge weights = **similarities** (e.g., Gaussian kernel)
- Matrix form: **adjacency matrix  $W$**

The choice of similarity function and neighborhood size strongly influences results.

# Graph Laplacian and Embedding

Compute a Laplacian matrix of the graph, such as:

$$L = D - W \quad \text{or} \quad L_{\text{sym}} = I - D^{-1/2}WD^{-1/2}$$

where  $D$  is the [degree matrix](#).

Extract the first  $k$  [eigenvectors](#) of  $L$  to form an embedding in  $\mathbb{R}^k$  that separates clusters.

# Clustering Step

Apply a standard clustering algorithm (typically **K-means**) to the rows of the eigenvector matrix.

- These rows represent the data in a **spectral embedding**.
- Cluster assignments in this space map back to clusters in the original data.

This combination captures structure missed by purely distance-based methods.

# Advantages and Limitations

## Advantages

- Captures **non-linear** and **arbitrary-shaped** clusters.
- Works well when clusters are connected components in a graph.

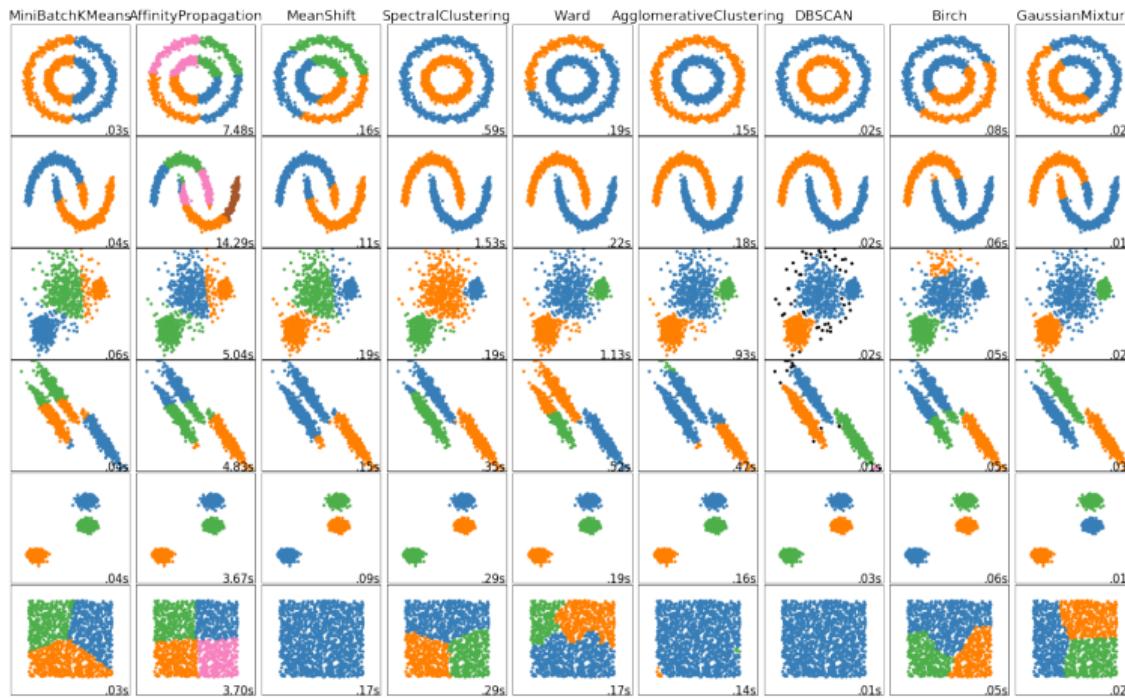
## Limitations

- Requires constructing and storing a similarity matrix.
- Not ideal for **very large** datasets due to eigenvector computation.

|   |                          |    |
|---|--------------------------|----|
| 1 | Hierarchical clustering  | 2  |
| 2 | Density based clustering | 27 |
| 3 | Model based clustering   | 47 |
| 4 | Final remarks            | 65 |

# Comparison of results for selected algorithms

(from Scikit-Learn documentation)



# A summary of selected clustering algorithms – I

| Method name                  | Parameters                               | Scalability                                                 | Use case                                                                            | Geometry (metric used)                       |
|------------------------------|------------------------------------------|-------------------------------------------------------------|-------------------------------------------------------------------------------------|----------------------------------------------|
| K-Means                      | number of clusters                       | Very large n_samples, medium n_clusters with MiniBatch code | General-purpose, even cluster size, flat geometry, not too many clusters, inductive | Distances between points                     |
| Affinity propagation         | damping, sample preference               | Not scalable with n_samples                                 | Many clusters, uneven cluster size, non-flat geometry, inductive                    | Graph distance (e.g. nearest-neighbor graph) |
| Mean-shift                   | bandwidth                                | Not scalable with n_samples                                 | Many clusters, uneven cluster size, non-flat geometry, inductive                    | Distances between points                     |
| Spectral clustering          | number of clusters                       | Medium n_samples, small n_clusters                          | Few clusters, even cluster size, non-flat geometry, transductive                    | Graph distance (e.g. nearest-neighbor graph) |
| Ward hierarchical clustering | number of clusters or distance threshold | Large n_samples and n_clusters                              | Many clusters, possibly connectivity constraints, transductive                      | Distances between points                     |

# A summary of selected clustering algorithms – II

| Method name              | Parameters                                                       | Scalability                             | Use case                                                                                         | Geometry (metric used)            |
|--------------------------|------------------------------------------------------------------|-----------------------------------------|--------------------------------------------------------------------------------------------------|-----------------------------------|
| Agglomerative clustering | number of clusters or distance threshold, linkage type, distance | Large n_samples and n_clusters          | Many clusters, possibly connectivity constraints, non Euclidean distances, transductive          | Any pairwise distance             |
| DBSCAN                   | neighborhood size                                                | Very large n_samples, medium n_clusters | Non-flat geometry, uneven cluster sizes, outlier removal, transductive                           | Distances between nearest points  |
| OPTICS                   | minimum cluster membership                                       | Very large n_samples, large n_clusters  | Non-flat geometry, uneven cluster sizes, variable cluster density, outlier removal, transductive | Distances between points          |
| Gaussian mixtures        | many                                                             | Not scalable                            | Flat geometry, good for density estimation, inductive                                            | Mahalanobis distances to centers  |
| BIRCH                    | branching factor, threshold, optional global clusterer.          | Large n_clusters and n_samples          | Large dataset, outlier removal, data reduction, inductive                                        | Euclidean distance between points |

# Clustering types

- Partitioning
  - iteratively find partitions in the dataset, optimizing some quality criterion
- Hierarchic
  - recursively compute a structured hierarchy of subsets
- Density based
  - compute densities and aggregates clusters in high density areas
- Model based
  - assume a model for the distribution of the data and find the model parameters which guarantee the best fitting to the data

# Clustering scalability

- Effectiveness decreases with
  - dimensionality  $D$
  - noise level
- Computational cost increases with
  - dataset size  $N$ , at least linearly
  - dimensionality  $D$

# Research perspective

- From the past
  - well-known problem in statistics
  - recent research
    - machine learning
    - databases
    - visualization
- For the future
  - effective and efficient algorithms for big data clustering, with a large number of dimensions, noise requested scalability w.r.t.:
    - size ( $N$ )
    - dimensionality ( $D$ )
    - noise level
    - rate of new data arrivals

# Uses of clustering – data comprehension

- Biology
  - Creation of taxonomies
  - Genetics
- Information Retrieval
  - Grouping documents
- Climatology
  - Repetition patterns
- Psychology and medicine
  - Identification of illness types in front of partial variation of evidence
- Business
  - Customer grouping

# Uses of clustering – utilities

- Summarization
  - Reasoning with groups representatives instead of with the entire population
- Data compression
  - Reduce the amount of data
  - Find cluster prototypes and substitute data with the indexes of the prototypes
    - Vector quantization
- Find the nearest neighbours
  - Each object refers to his prototypes
  - Near neighbours refer to the same prototype

# Bibliography I

- ▶ Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu.  
A density-based algorithm for discovering clusters in large spatial databases with noise.  
pages 226–231. AAAI Press, 1996.
- ▶ Pang-Nin Tan, Michael Steinbach, and Vipin Kumar.  
[Data Mining](#).  
Addison Wesley, 2006.  
ISBN 0-321-32136-7.