

Machine Learning

Clustering

Claudio Sartori

DISI

Department of Computer Science and Engineering – University of Bologna, Italy

claudio.sartori@unibo.it

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
5	Density based clustering	94
6	Model based clustering	114
7	Final remarks	123

The problem of clustering

For a comprehensive review see [Jain et al.(1999) Jain, Murty, and Flynn]

- **Given** – a set of N objects x_i , each described by D values x_{id}
- **Task** – find a *natural* partitioning in K clusters and, possibly, a number of *noise* objects
- **Result** – a *clustering scheme*, i.e. a function mapping each data object to the sequence $[1 \dots K]$ (or to noise)

Desired property of clusters

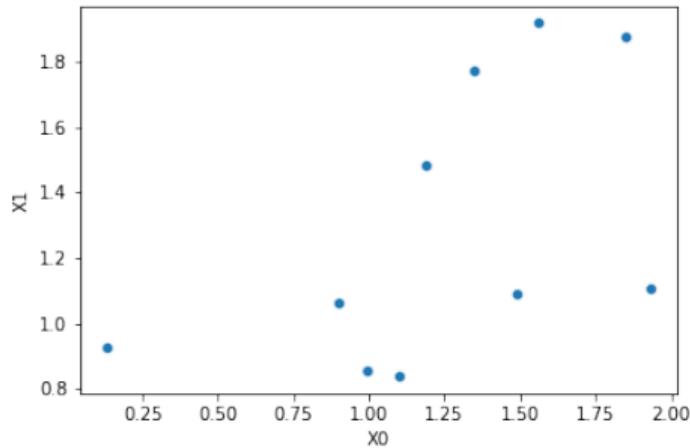
- **objects in the same cluster are similar**
 - look for a clustering scheme which maximizes intra-cluster similarity

A little bit of formality

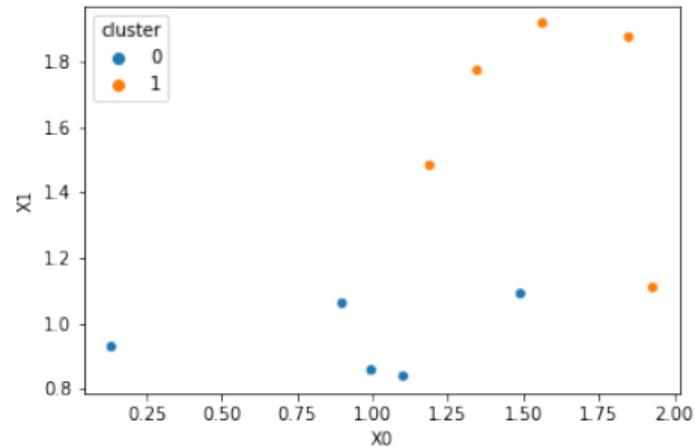
- find a function $clust()$ from \mathcal{X} to $1..K$ such that:
 - $\forall x_1, x_2 \in \mathcal{X}$, $clust(x_1) = clust(x_2)$ when x_1 and x_2 are *similar*
 - $\forall x_1, x_2 \in \mathcal{X}$, $clust(x_1) \neq clust(x_2)$ when x_1 and x_2 are *not similar*

Clustering

2-d Data



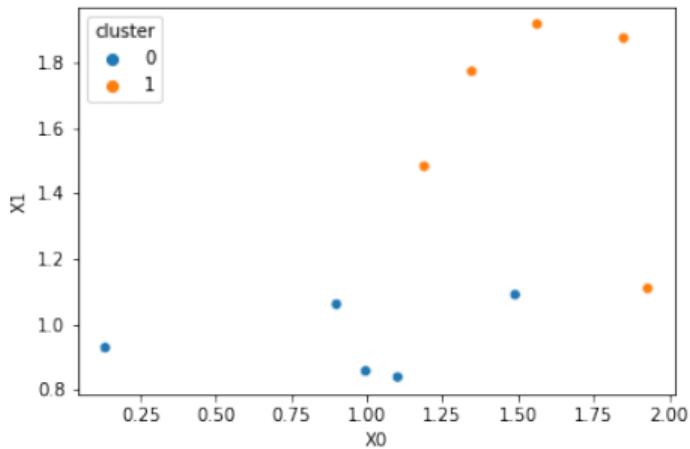
2-d Data clustered



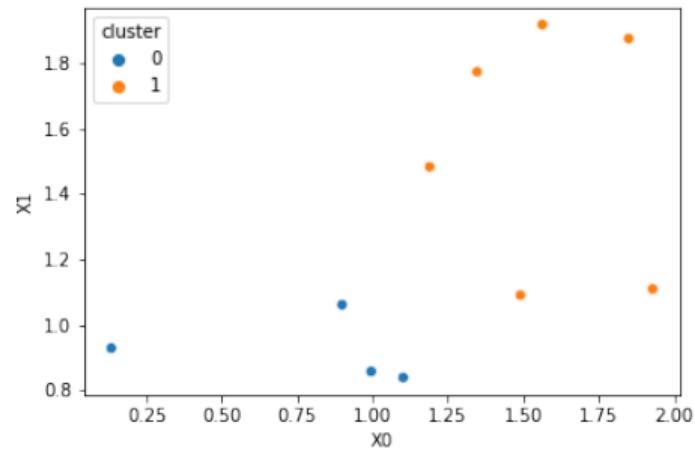
The clustering function maps points to clusters

Clustering - two different clustering functions

$clust^a$



$clust^b$



Which of the two is better, i.e. maximises intra-cluster similarity?
A measure is needed

Ideas for a measure?

- the sum of the distances between a point and the others in the same cluster?
 -
- the average of the distances between a point and the others in the same cluster?
 -
- the sum of the squared distances?
 -
- we could choose a point which, in some sense, **represents** the points of the cluster
 -
- ...
 -

Ideas for a measure? Discussion

- the sum of the distances between a point and the others in the same cluster?
 - bigger for bigger clusters
- the average of the distances between a point and the others in the same cluster?
 - not influenced by cluster size
- the sum of the squared distances?
 - more penalisation for **sparsity**
- we could choose a point which, in some sense, **represents** the points of the cluster
 - what about the **average of the coordinates?**
- ...

Centroid

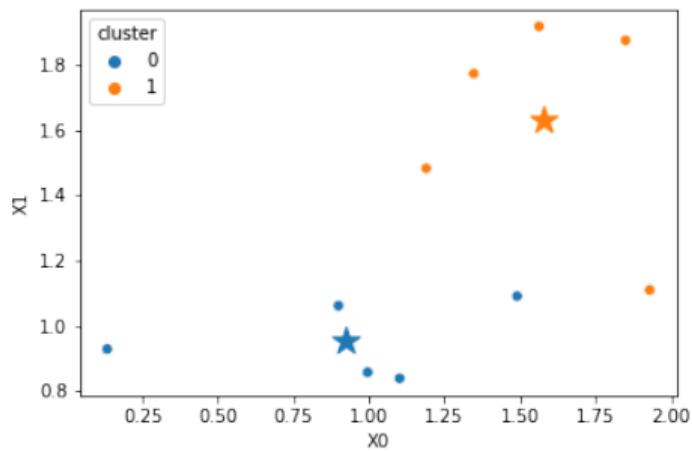
- a point with coordinates computed as the average of the coordinates of all the points in the cluster
- in physics it is the **center of gravity** of a set of points of equal mass
- for each cluster k and dimension d , the d coordinate of the **centroid** is

$$\text{centroid}_d^k = \frac{1}{|x_i : \text{clust}(x_i) = k|} \sum_{x_i : \text{clust}(x_i) = k} x_{id}$$

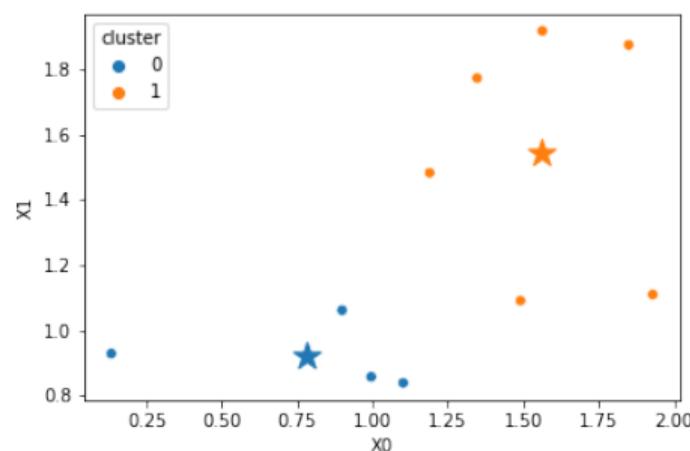
Clustering - two different clustering functions

The stars represent the *centroids*

$clust^a$



$clust^b$



Which of the two is better, i.e. maximises intra-cluster similarity?
A measure is needed

Taxonomy of the clustering methods

- Partitioning
 - K-means (MacQueen 67), expectation maximization (Lauritzen 95), CLARANS (Ng and Han 94)

- Hierarchic
 - agglomerative/divisive, BIRCH (Zhang et al 96), CURE (Guha et al 98)

- Based on linkage
- Based on density

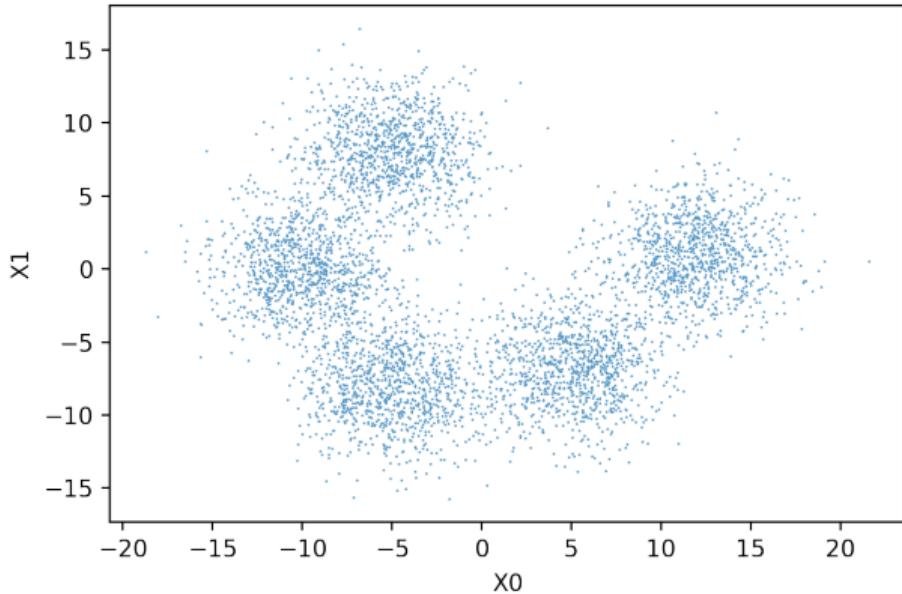
- DBSCAN (Ester et al 96), DENCLUE (Hinnenburg and Keim 98)

- Statistics
 - IBM-IM demographic clustering, COBWEB (Fisher 87), Autoclass (Cheeseman 96)

1	Introduction to clustering	2
2	K-means	12
	● Minimize distortion	29
	● Issues about K-means	35
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
5	Density based clustering	94
6	Model based clustering	114
7	Final remarks	123

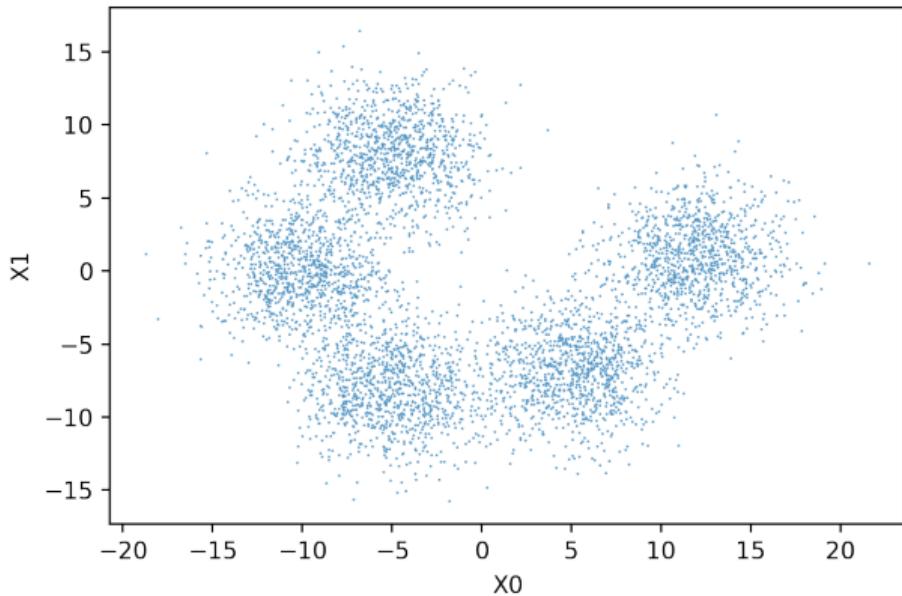
Some data

- Could be modeled as a five component gaussian mixture
 - ... *statistician voice*...
- How do we guess the number five in a D -dimensional space (with $D \geq 2$)?



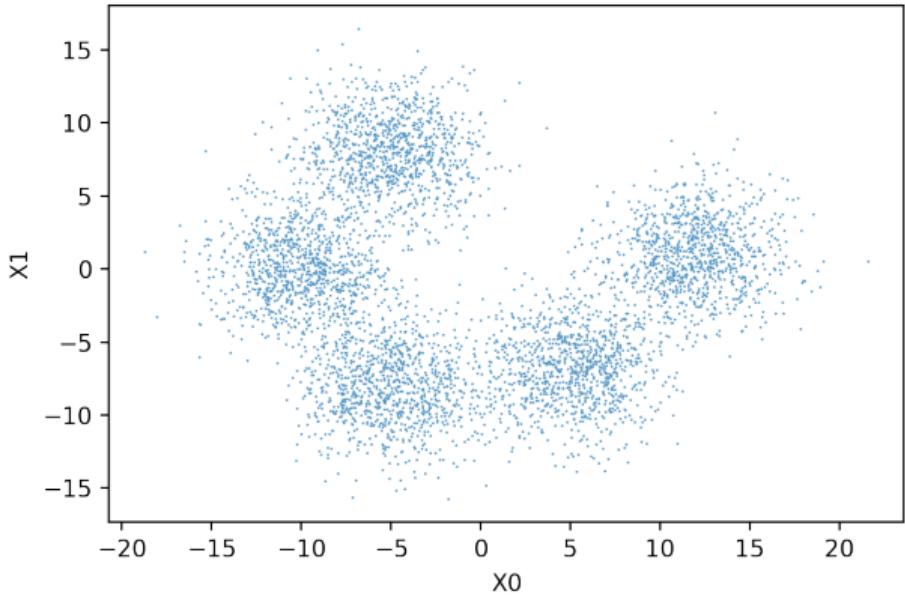
Transmission

- Transmit the coordinates of points
- Allow only two bits per point
 - the transmission will be **lossy**
- Need a coding/decoding mechanism



How much loss?

- Sum of the squared errors between the real points and their encoding/decoding
- Which encoding/decoding minimizes the loss?

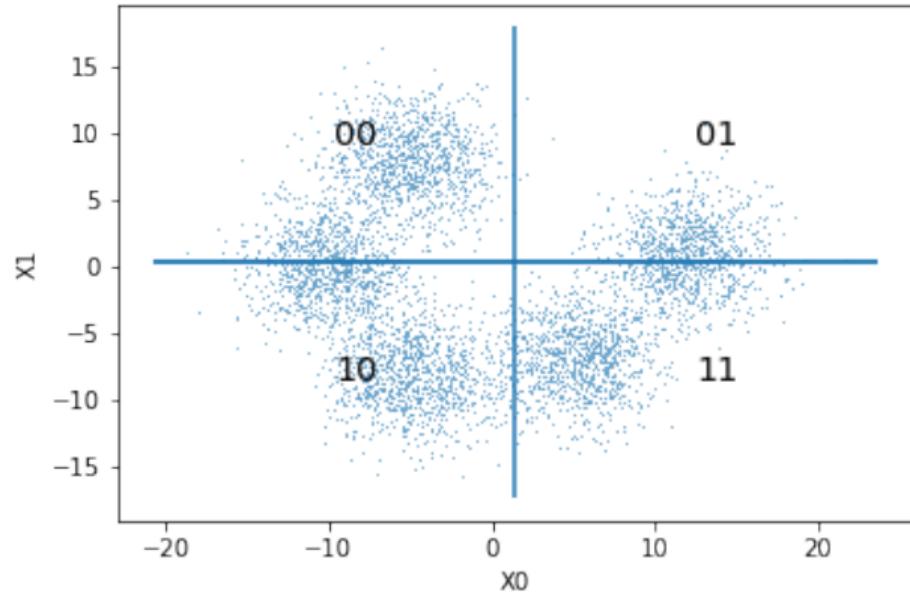


How much loss?

- Sum of the squared errors between the real points and their encoding/decoding

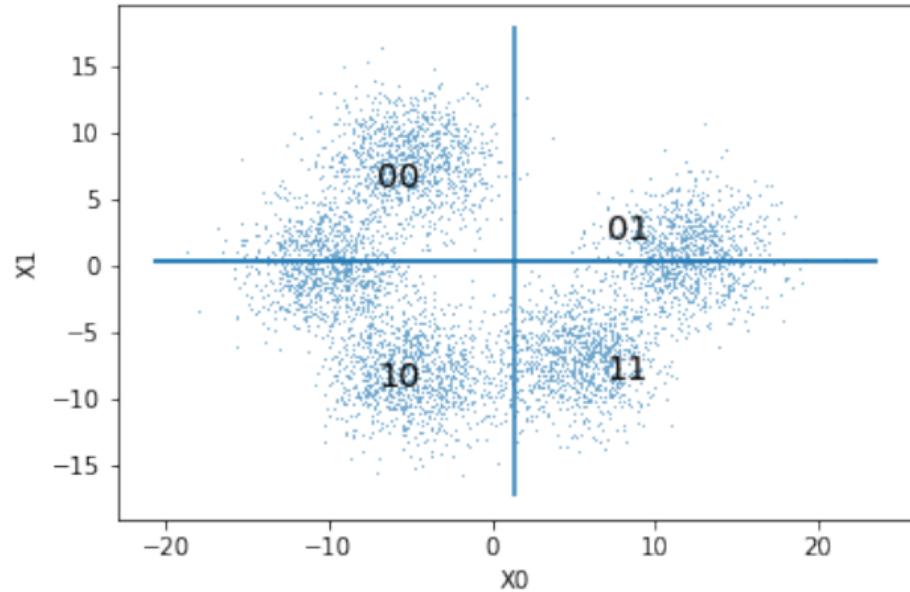
First idea

- partition the space into a grid of cells
- decode each pair of bits with the center of the grid cell



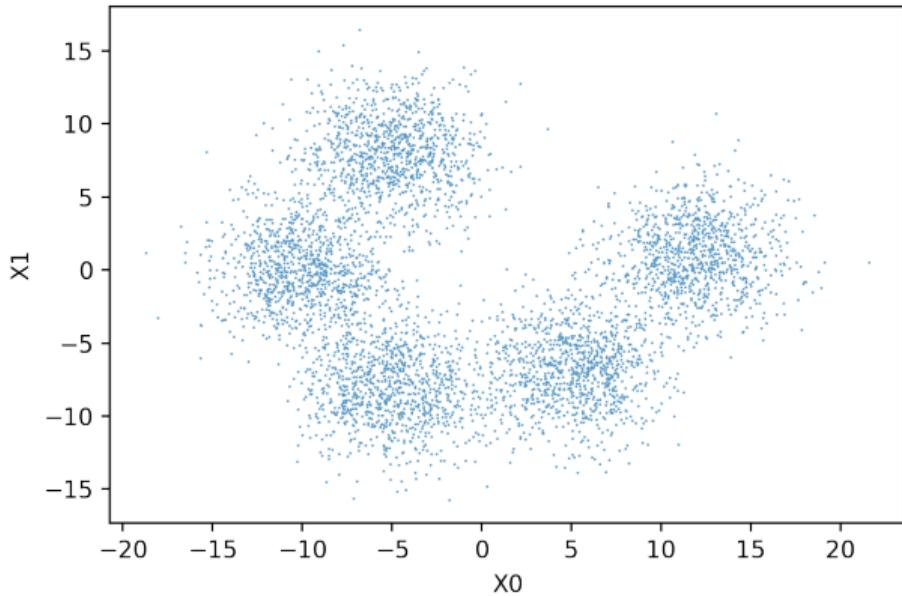
Improvement

- partition the space into a grid of cells
- decode each pair of bits with the centroid of the points in the grid cell



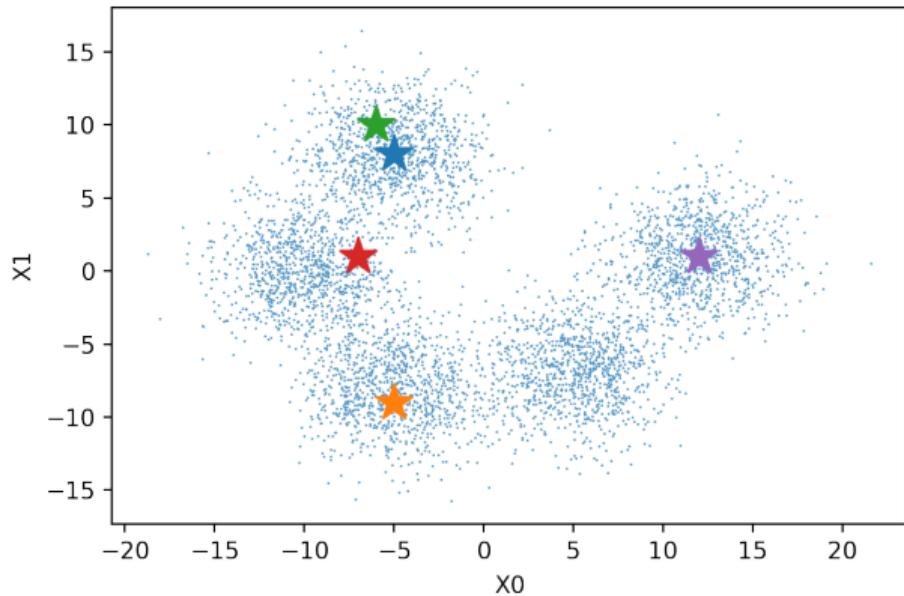
K-means

1. Ask the user the number of clusters K
1.1 $\Rightarrow 5$



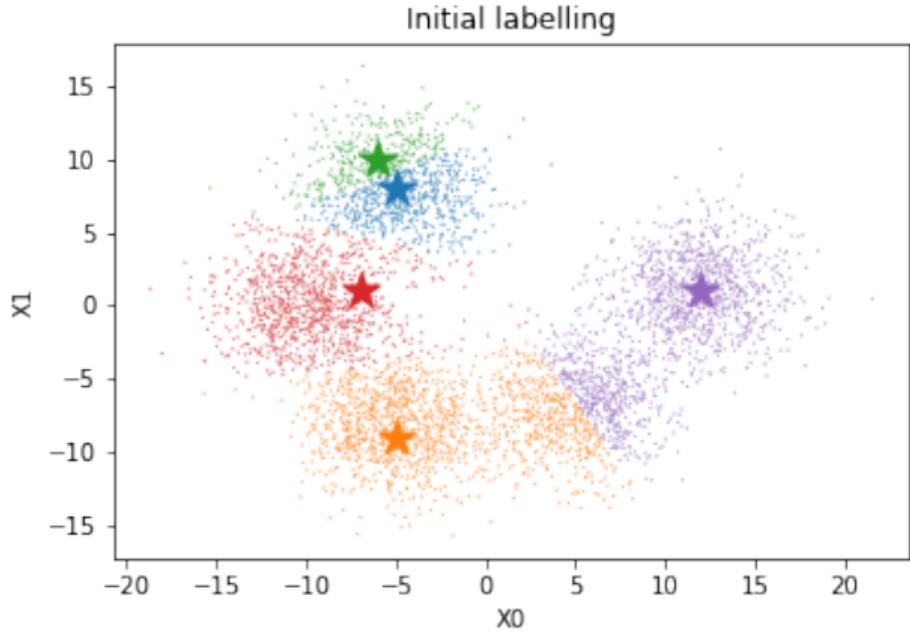
K-means

1. Ask the user the number of clusters K
2. Random choice of K points as **temporary centers**



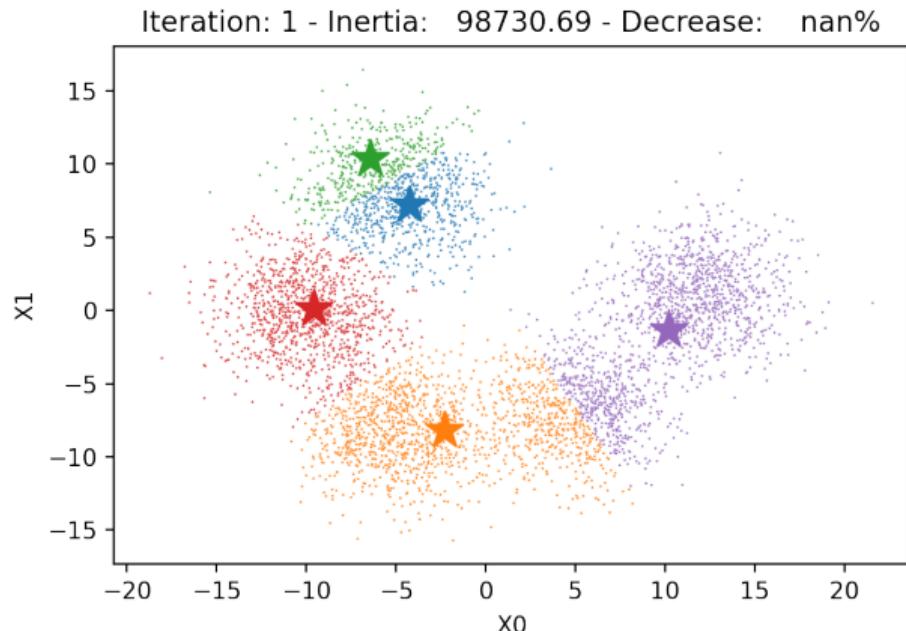
K-means

1. Ask the user the number of clusters K
2. Random choice of K points as **temporary centers**
3. Each point finds his nearest center and is labelled (i.e. colored) accordingly

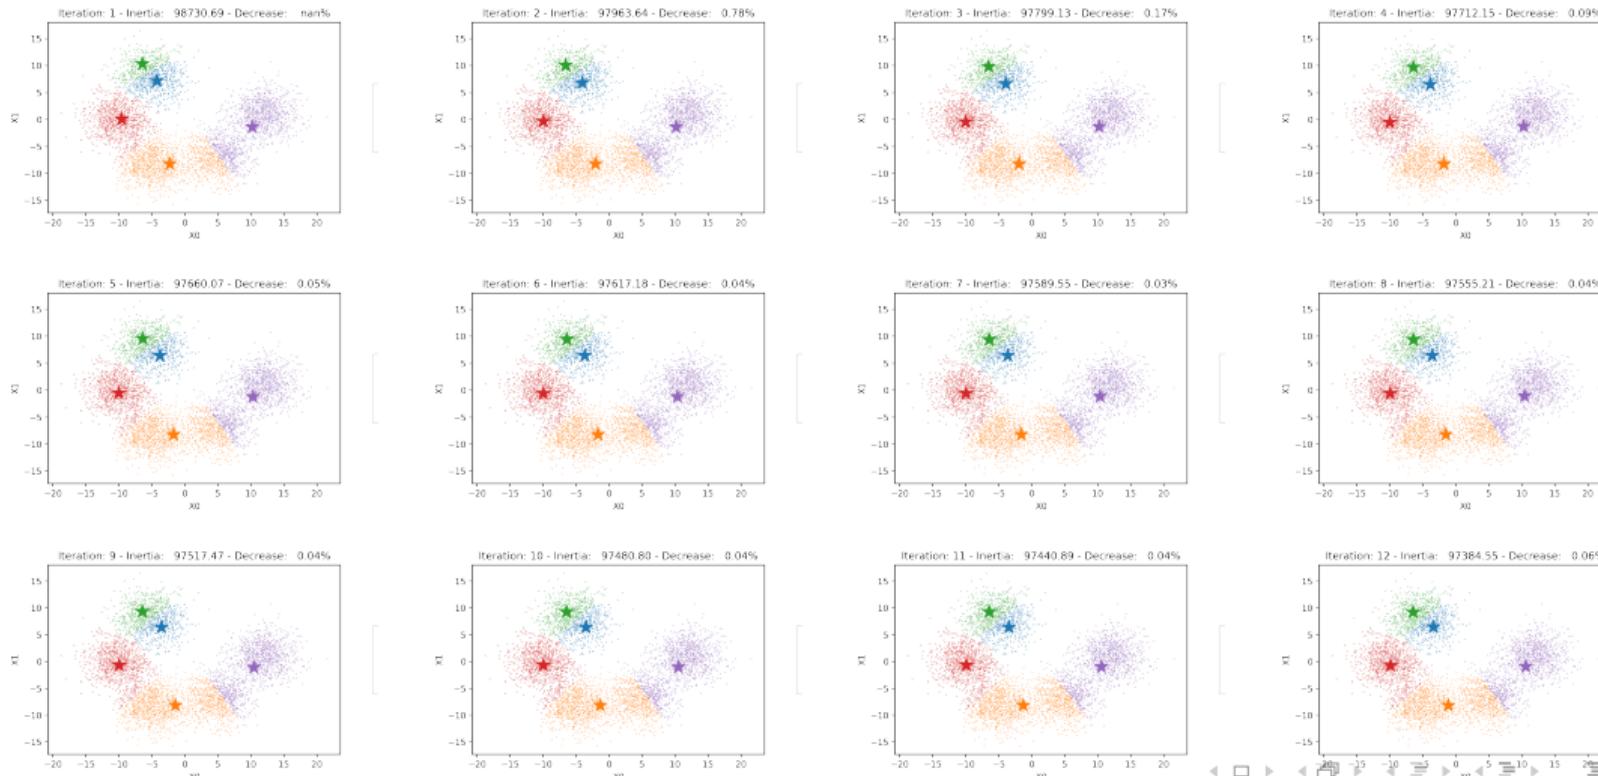


K-means

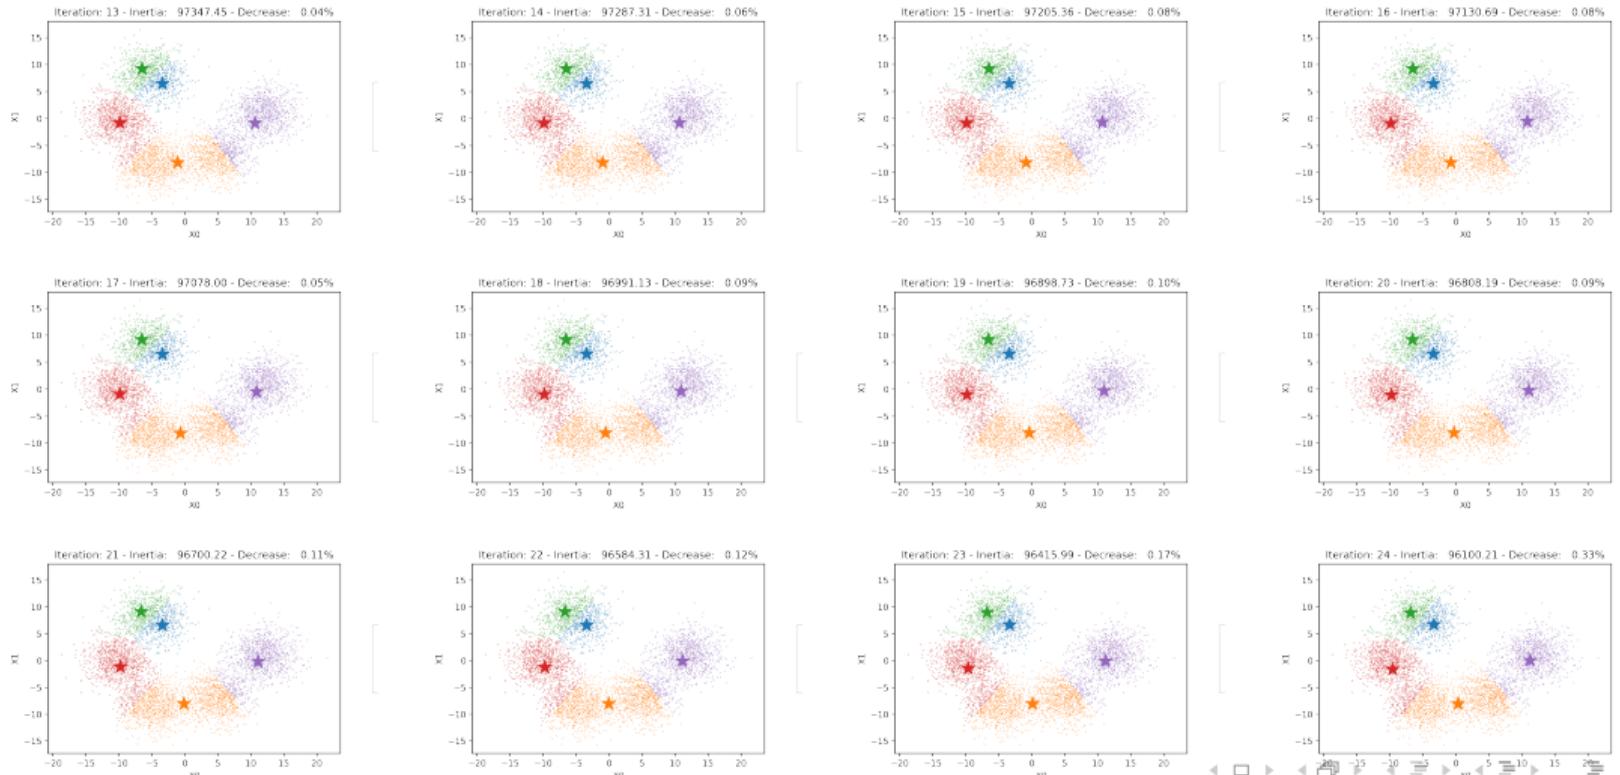
1. Ask the user the number of clusters K
2. Random choice of K points as **temporary centers**
3. Each point finds his nearest center
4. for each center finds the centroid of its points ...
5. ... and move there the center



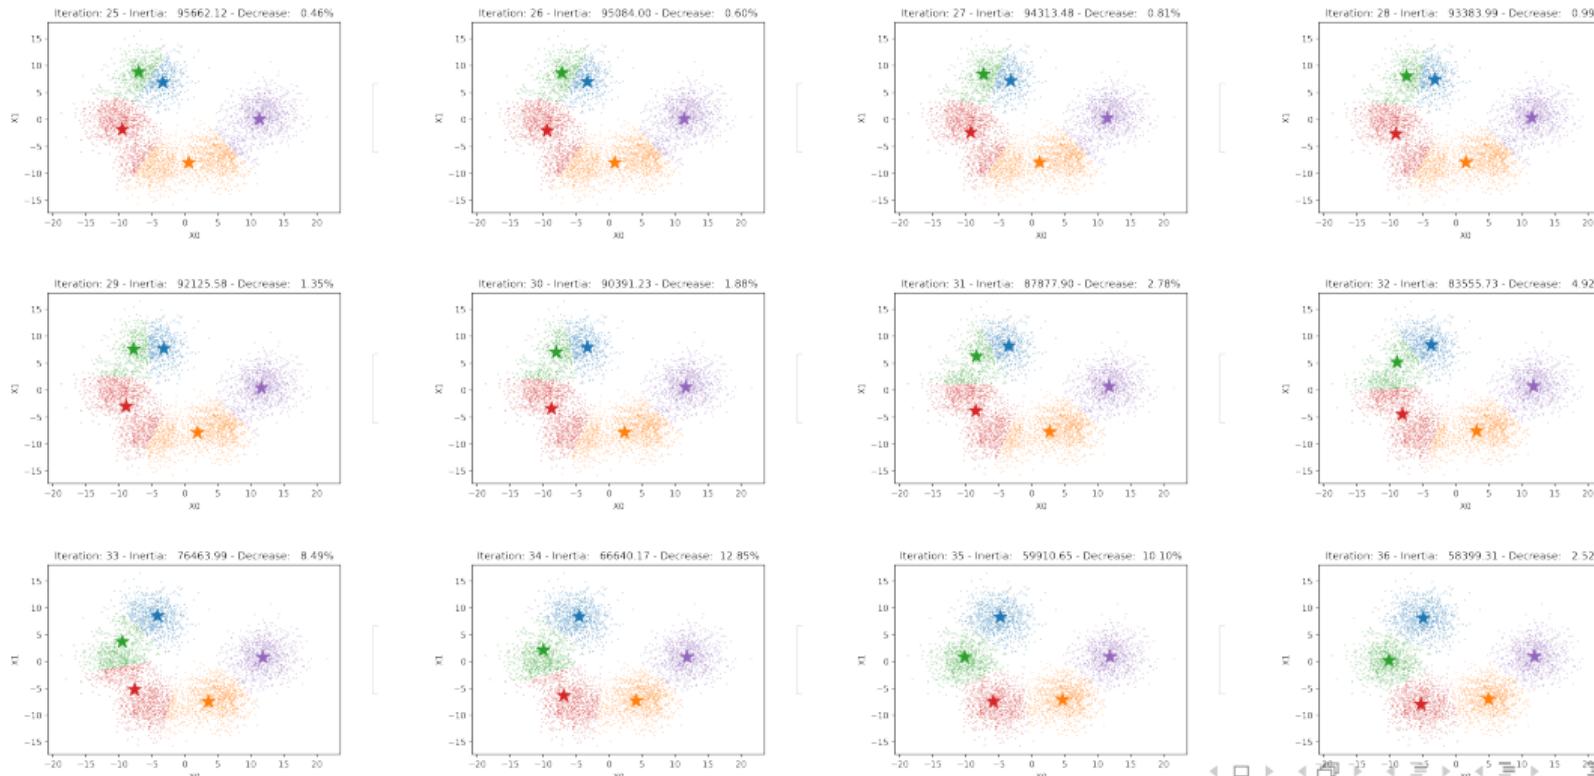
KMeans working . . .



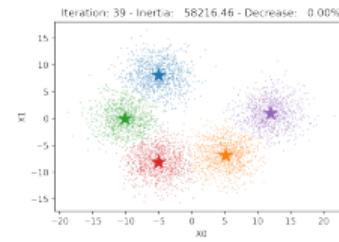
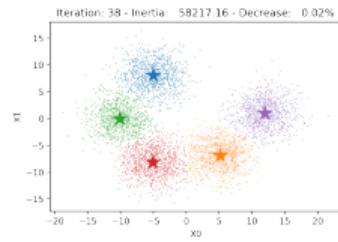
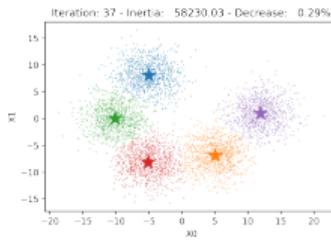
KMeans working . . .



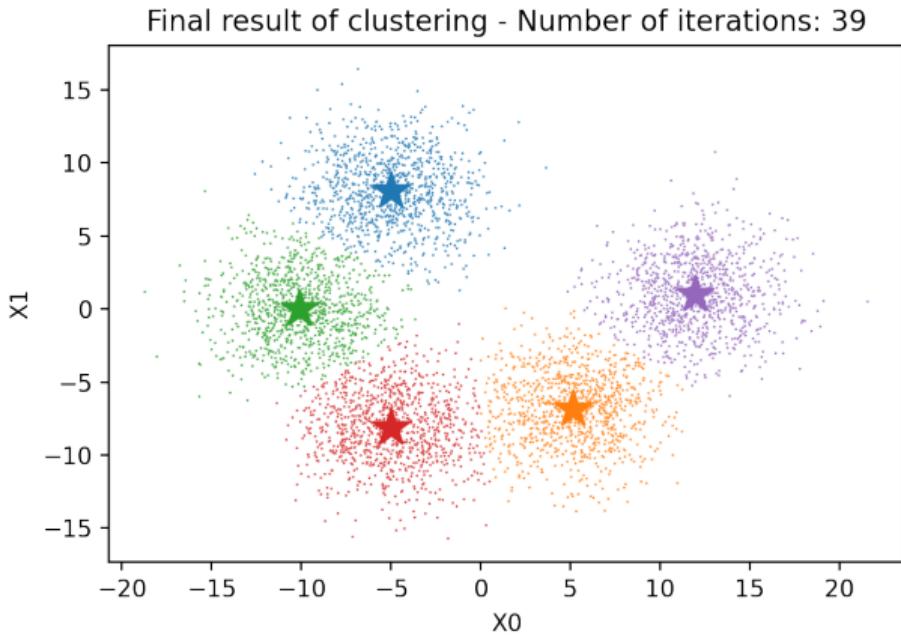
KMeans working . . .



KMeans working . . .



K-means ends



Questions

1. What are we trying to optimize?
2. Is termination guaranteed?
3. Are we sure that the best clustering scheme is found?
 - 3.1 Which is the definition of best clustering scheme?
4. How should we start?
5. How can we find the number of clusters?

Question 1: Distortion

Frequently called in the literature **Inertia**

Given:

- a dataset $\{x_i, i = 1 \dots N\}$
- a coding function $\text{Encode} : \{R\}^D \rightarrow [1..K]$
- a decoding function $\text{Decode} : [1..K] \rightarrow \mathbb{R}^D$
- define $\text{Distortion} = \sum_{i=1}^N (x_i - \text{Decode}(\text{Encode}(x_i)))^2$
- shortcut $\text{Decode}(k) = \mathbf{c}_k$

then

$$\text{Distortion} = \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2$$

Minimal distortion I

$$\text{Distortion} = \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2$$

Which properties are requested to $\mathbf{c}_1, \dots, \mathbf{c}_K$ for the minimal distortion?

1. x_i must be encoded with the nearest center

Why?

Because otherwise the distortion could be reduced by substituting $\text{Encode}(x_i)$ with the nearest center

$$\mathbf{c}_{\text{Encode}(x_i)} = \operatorname*{argmin}_{\mathbf{c}_j \in \{\mathbf{c}_1, \dots, \mathbf{c}_K\}} (x_i - \mathbf{c}_j)^2$$

Minimal distortion II

$$\text{Distortion} = \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2$$

Which properties are requested to $\mathbf{c}_1, \dots, \mathbf{c}_K$ for the minimal distortion?

2. The partial derivative of distortion w.r.t. the position of each center must be zero

Why?

Because in that case the function has either a maximum or a minimum

Step 2. The partial derivative of distortion w.r.t. the position of each center must be zero

$$\begin{aligned}\text{Distortion} &= \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2 \\ &= \sum_{j=1}^K \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} (x_i - \mathbf{c}_j)^2\end{aligned}$$

$$\begin{aligned}\frac{\partial \text{Distortion}}{\partial \mathbf{c}_j} &= \frac{\partial}{\partial \mathbf{c}_j} \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} (x_i - \mathbf{c}_j)^2 \\ &= -2 \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} (x_i - \mathbf{c}_j)\end{aligned}$$

2. The partial derivative of distortion w.r.t. the position of each center must be zero

When distortion is minimal

$$\mathbf{c}_j = \frac{1}{|\text{OwnedBy}(\mathbf{c}_j)|} \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} x_i$$

Minimal distortion III

$$\text{Distortion} = \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2$$

Which properties are requested to $\mathbf{c}_1, \dots, \mathbf{c}_K$ for the minimal distortion?

1. x_i must be encoded with the nearest center
2. each center must be the **centroid** of the points it owns

Algorithm: Improving a sub-optimal solution

$$\text{Distortion} = \sum_{i=1}^N (x_i - \mathbf{c}_{\text{Encode}(x_i)})^2$$

Which properties are requested to $\mathbf{c}_1, \dots, \mathbf{c}_K$ for the minimal distortion?

1. x_i must be encoded with the nearest center
2. each center must be the **centroid** of the points it owns

⇒ Alternately perform steps 1 and 2

It can be proven that after a finite number of steps the system reaches a state where neither of the two operations changes the state

Why?

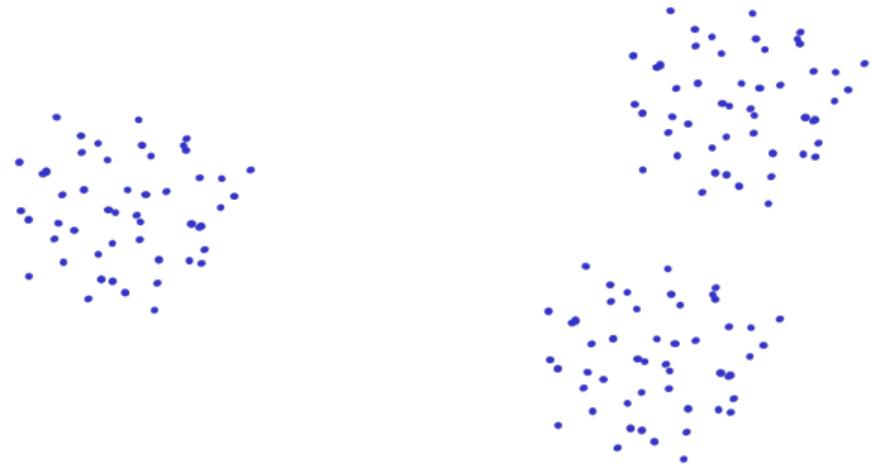
Question 2: Algorithm termination

- There is only a finite number of ways to partition N objects into K groups
- The state of the algorithm is given by the two encode/decode functions
- The number of configurations where all the centers are the centroids of the points they own is **finite**
- If after one iteration the state changes, the distortion is **reduced**
- Therefore each change of state bring to a state which was never visited before
- In summary, sooner or later the algorithm will stop because there are no new states reachable

Question 3: Local or global minimum?

Is the ending state the best possible?

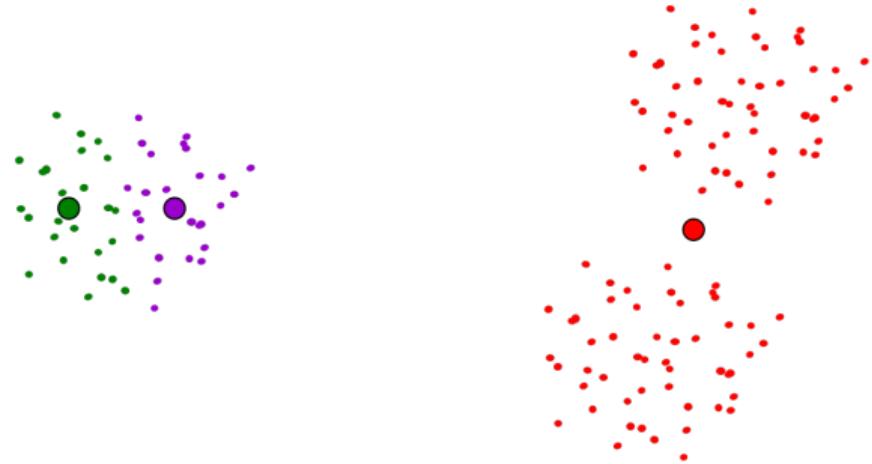
- Not necessarily
- An example



Question 3: Local or global minimum?

Is the ending state the best possible?

- Not necessarily
- An example



Question 4: Looking for a good ending state

- The starting point is important
 - choose randomly the first starting point
 - choose in sequence the $2..K$ starting points as far as possible from the preceding ones
- Re-run the algorithm with different starting points

Question 5: Choose the number of clusters

not so easy...

- try various values
- use a **quantitative evaluation** of the quality of the clustering scheme to decide among the different values
- the best value finds the optimal compromise between the minimization of intra-cluster distances and the maximization of the inter cluster distances

The proximity function

- The most obvious solution, used in the previous formulas is the euclidean distance
 - good choice, in general, for vector spaces
- Several alternative solutions for specific data types and data sets
 - see the "Data" module for additional discussions

Sum of Squared Errors I

The official name of the distortion

$$\begin{aligned} \text{SSE} &= \sum_{j=1}^K \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} (x_i - \mathbf{c}_j)^2 \\ &= \sum_{j=1}^K \text{SSE}_j \end{aligned}$$

Sum of Squared Errors II

- A cluster j with high SSE_j has low quality
- $SSE_j = 0$ if and only if all the points are coincident with the centroid
- SSE decreases for increasing K , is zero when $K = N$
- \Rightarrow minimizing SSE is not a viable solution to choose the best K
 - more discussions on this in the section on “Evaluation of the quality of a clustering scheme”

Outliers

- are points with high distance from their centroid
 - high contribution to SSE
- have a bad influence on the clustering results
 - sometimes it is a good idea to remove them
 - the choice is related to the application domain

Common uses of K-means

- It can be easily used in the beginning, for the exploration of data
- In a one-dimension space it is a good way to discretize the values of a domain in non-uniform buckets
- It is the basis for vector quantization, a classical technique for signal processing and compression
- Used for choosing the color palettes
 - gif compressed images: color quantization

Complexity

Given:

- T number of iterations
- K number of clusters
- N number of data points
- D number of dimensions

the time complexity is

$$\mathcal{O}(TKND)$$

Pros and cons of K-means

- Strong points
 - fairly efficient, nearly linear in the number of data points
 - in general $T, K, D \ll N$
- Weak points
 - in essence it is defined for spaces where the centroid can be computed
 - e.g. when the Euclidean distance is available, also other distance functions work well
 - cannot work with nominal data
 - requires the K parameter
 - nevertheless the best K can be found with iterations
 - it is very sensitive to outliers
 - does not deal with **noise**
 - does not deal properly with **non convex** clusters

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
	● Cohesion and separation	51
	● Silhouette	55
	● Choice of K	60
	● Supervised measures	64
4	Hierarchical clustering	69
5	Density based clustering	94
6	Model based clustering	114
7	Final remarks	123

Evaluation of a clustering scheme

- It is related only to the result, not to the clustering technique
- Clustering is a non supervised method
 - the evaluation is critical, because whenever there is very little apriori information, such as class labels
 - we need one or more **score** function to measure various properties of the clusters and of the clustering scheme as a whole
 - in the literature the words *score* and *index* are considered synonyms in this context
 - if some supervised data are available, they can be used to evaluate the clustering scheme
- In 2D the clusters can be examined visually
- In higher order spaces the 2D projections can help, but in general it is better to use more formal methods

Issues on the evaluation of clustering

- Distinguish patterns from random apparent regularities
- Find the best number of clusters
- Non supervised evaluation
- Supervised evaluation
- Relative comparison of clustering schemes

Proximity and others

- Similarity – Proximity
 - a two variable function measuring how much two objects are **similar**, according to the values of their properties
- Dissimilarity
 - a two variable function measuring how much two objects are **different**, according to the values of their properties
 - e.g. the **Euclidean distance**

Global separation of a clustering scheme

SSB – Sum of Squares Between clusters

\mathbf{c} = global centroid of the dataset

$$\text{SSB} = \sum_{i=1}^K N_i \text{Dist}(\mathbf{c}_i, \mathbf{c})^2$$

Link between cohesion and separation – I

- TSS = Total Sum of Squares
 - sum of squared distances of the points from the global centroid
- $TSS = SSE + SSB$
- the total sum of squares is a global property of the dataset, independent from the clustering scheme
- for a given dataset, minimise SSE \Leftrightarrow maximise SSB

Link between cohesion and separation – II

$$\begin{aligned} \text{TSS} &= \sum_{i=1}^K \sum_{x \in k_i} (x - \mathbf{c})^2 = \sum_{i=1}^K \sum_{x \in k_i} ((x - \mathbf{c}_i) - (\mathbf{c} - \mathbf{c}_i))^2 \\ &= \sum_{i=1}^K \sum_{x \in k_i} (x - \mathbf{c}_i)^2 - 2 \sum_{i=1}^K \sum_{x \in k_i} (x - \mathbf{c}_i)(\mathbf{c} - \mathbf{c}_i) + \sum_{i=1}^K \sum_{x \in k_i} (\mathbf{c} - \mathbf{c}_i)^2 \\ &= \sum_{i=1}^K \sum_{x \in k_i} (x - \mathbf{c}_i)^2 + \sum_{i=1}^K \sum_{x \in k_i} (\mathbf{c} - \mathbf{c}_i)^2 \\ &= \sum_{i=1}^K \sum_{x \in k_i} (x - \mathbf{c}_i)^2 + \sum_{i=1}^K |k_i|(\mathbf{c} - \mathbf{c}_i)^2 = \text{SSE} + \text{SSB} \end{aligned}$$

since $\sum_{x \in k_i} (x - \mathbf{c}_i) = 0$ by definition of \mathbf{c}_i

Evaluation of specific clusters and objects

- Each cluster can have its own evaluation
 - the worst clusters can be considered for additional split
- A weakly separated pair of clusters could be considered for merging
- Single objects can give negative contribution to the cohesion of a cluster or to the separation between two clusters
 - border objects

Silhouette score of a cluster – I

Requirements for a clustering quality score

- values are in a standard range, e.g. $-1, 1$
- increases with the separation between clusters
- decreases for clusters with low *cohesion*

$$\begin{aligned} \text{SSE} &= \sum_{j=1}^K \sum_{i \in \text{OwnedBy}(\mathbf{c}_j)} (x_i - \mathbf{c}_j)^2 \\ &= \sum_{j=1}^K \text{SSE}_j \end{aligned}$$

Silhouette score of a cluster – I

Requirements for a clustering quality score

- values are in a standard range, e.g. $-1, 1$
- increases with the separation between clusters
- decreases for clusters with low *cohesion*, or, in other words, with high *sparsity*

Silhouette score of a cluster – II

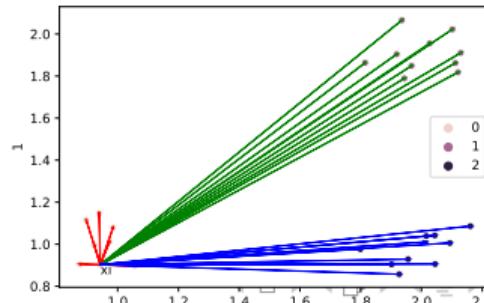
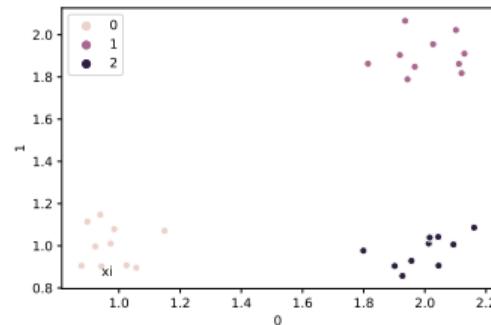
Consider the individual contribution of each object, say x_i

Contribution to cluster **sparsity**: the average of red distances

$$a_i = \text{average } \underset{j, y(x_j) = y(x_i)}{\text{dist}}(x_i, x_j)$$

Contribution to **separation** from other clusters: the minimum of the two averages of green and blue distances

$$b_i = \min_{k \in \mathcal{Y}, k \neq y(x_i)} (\text{average } \underset{j, y(x_j) = k}{\text{dist}}(x_i, x_j))$$



Silhouette score of a cluster – III

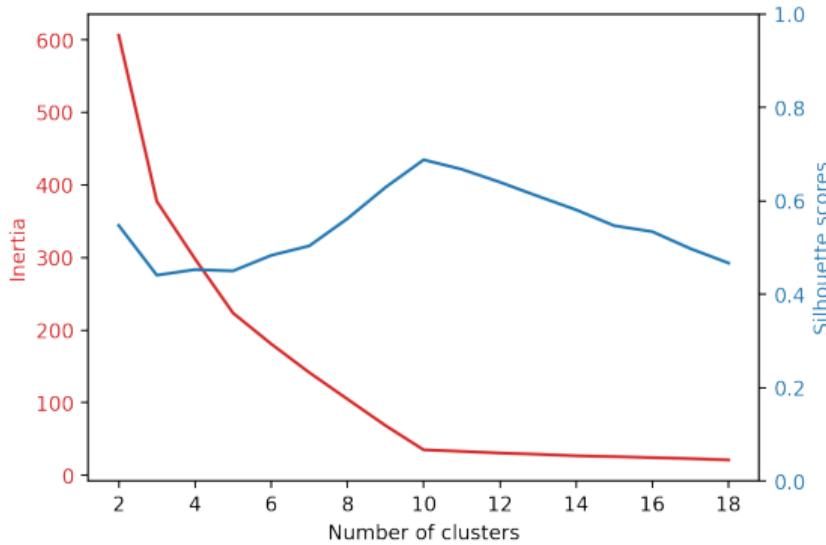
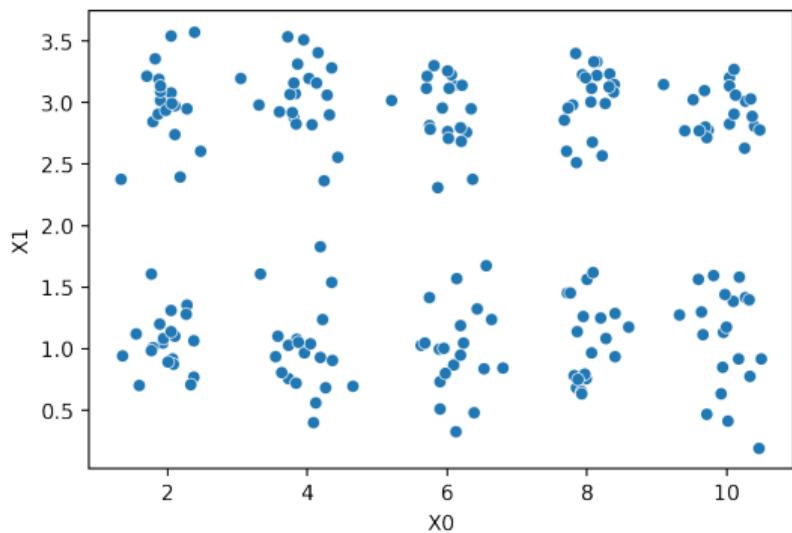
- Silhouette score of x_i

$$s_i = \frac{b_i - a_i}{\max(a_i, b_i)} \in [-1, 1]$$

- For the global score of a cluster/clustering scheme compute the average score over the cluster/dataset
- Intuition
 - when the score is less than zero for an object it means that there is a dominance of objects in other clusters at a distance smaller than objects of the same cluster

Example: Inertia and silhouette scores

Testing K-means with different numbers of clusters



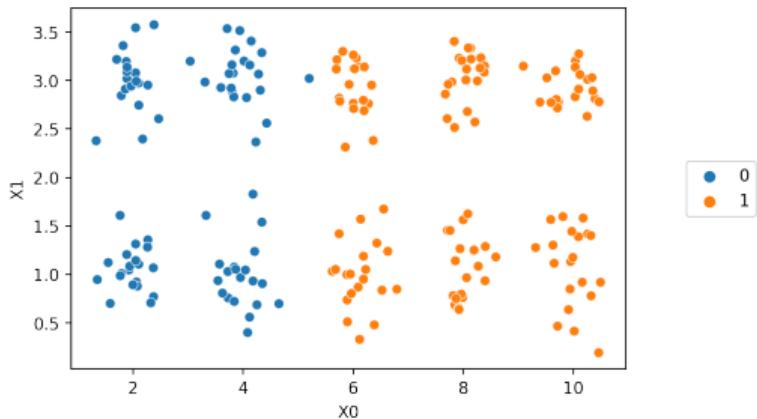
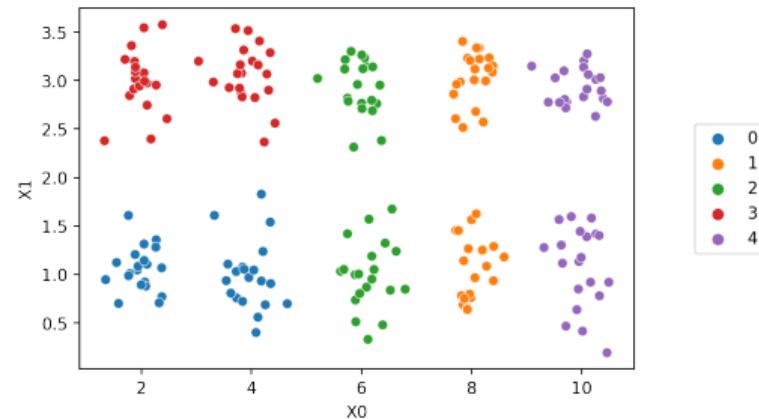
Looking for the best number of clusters – I

- Some algorithms, such as K-means, require the number of clusters as a parameter
- Measures, such as SSE and Silhouette, are obviously influenced by the number of clusters
 - they can be used to optimize K
- Computation of Silhouette score is expensive
- SSE decreases monotonically for increasing K
 - is equal to TSS for $K = 1$
 - goes to zero when $K = N$

Looking for the best number of clusters – II

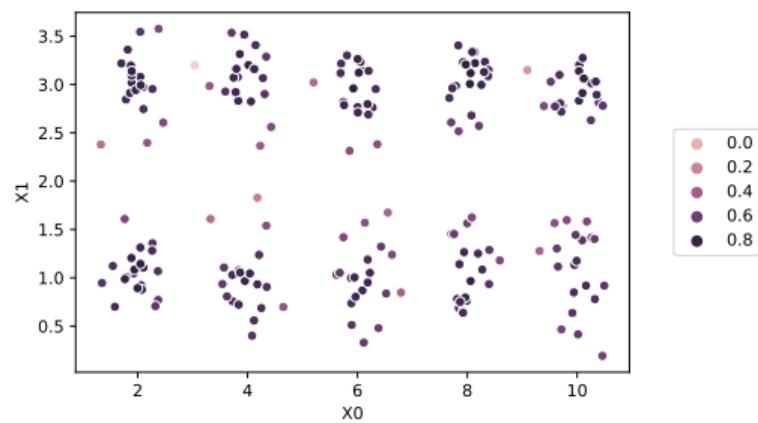
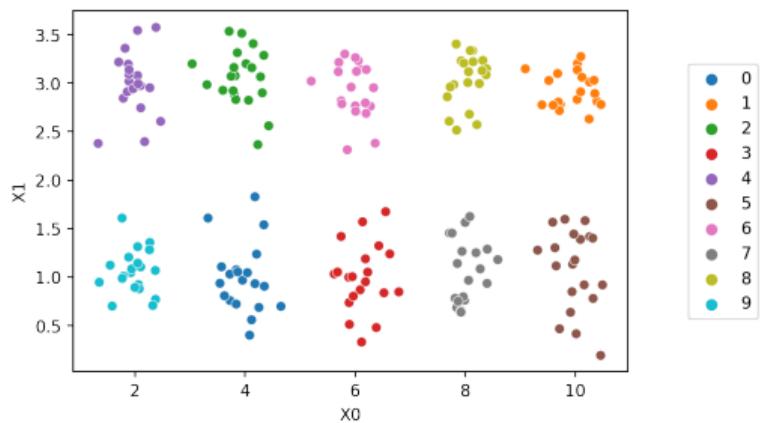
- The *inertia* varying K has frequently one or more points where the **slope decreases**: one of this points is frequently a plausible value for K
 - this is called **elbow method**
- The *silhouette* score varying K has frequently a maximum, in this case it indicates the best value for K

K-means results on the dataset of page 59

 $K = 2$  $K = 5$

K-means results on the dataset of page 59

Best *silhouette* for $K = 10$



Supervised measures: Gold Standard I

- Let be available a partition of a dataset similar to the data to be clustered, which we call **gold standard**, and defined by a labelling scheme $y_g(\cdot)$
 - it is the same as the labels attached to supervised data for training a classifier
- Consider a clustering scheme $y_k(\cdot)$
 - the cardinalities of the sets of distinct labels generated by the two schemes \mathcal{V}_g and \mathcal{V}_k can be different, and also in case of identity of the two grouping schemes, a permutation of labels could be necessary to make them equal

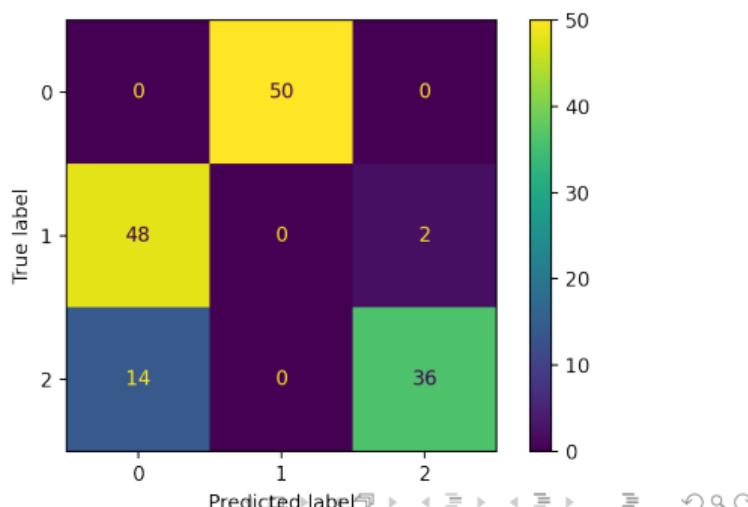
Why should we compare it with the gold standard?

- validate a clustering technique which can be applied later to new, unlabelled data
- the purpose is quite similar to testing a classifier
- the difference is that in this case we are more interested in *grouping* new data than in *labelling them* following the Gold Standard scheme

Classification-oriented measures

- Measure how the gold standard classes are distributed among the clusters
 - confusion matrix, precision, recall, f-measure
- On the right the confusion matrix for the **Iris** dataset
 - best match with permutation¹ of the predicted labels
 $0 \rightarrow 1, 1 \rightarrow 0, 2 \rightarrow 2$

```
X,y = load_iris(return_X_y=True)
estimator = KMeans(n_clusters=3
                    , random_state=363)
y_km = estimator.fit_predict(X)
disp = ConfusionMatrixDisplay(confusion_matrix(y,y_km))
disp.plot()
```



Similarity oriented measures - I

- Analogous to compare binary data
- Consider a clustering scheme $y_k(\cdot)$ and compare it with the **Gold Standard** $y_g(\cdot)$
- Any pair of objects can be labelled as
 - *SGSK* if they belong to the same set in $y_g(\cdot)$ and $y_k(\cdot)$
 - *SGDK* if they belong to the same set in $y_g(\cdot)$ and not in $y_k(\cdot)$
 - *DGSK* if they belong to the same set in $y_k(\cdot)$ but not in $y_g(\cdot)$
 - *DGDK* if they belong to different sets both in $y_g(\cdot)$ and $y_k(\cdot)$

Similarity oriented measures - II

Results given by
`pair_confusion_matrix(y_g,y_k)`

	SK	DK
SG	13512	1488
DG	1200	6150

Rand Score $\frac{SGSK + DGDK}{SGSK + DGDK + SGDK + DGSK} = 0.88$

Adjusted Rand Score Excludes the count of matches expected by chance¹ = 0.73

Jaccard Coefficient for label c $\frac{SG_c SK_c}{SG_c SK_c + SG_c DK_c + DG_c SK_c} = (1, 0.75, 0.69)$

- it requires remapping of $y_g(\cdot)$ to obtain the best match

1 See the [Wikipedia page](#) for a reference

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
	● Separation	72
	● Single linkage hierarchical clustering	74
	● Examples	82
5	Density based clustering	94
6	Model based clustering	114
7	Final remarks	123

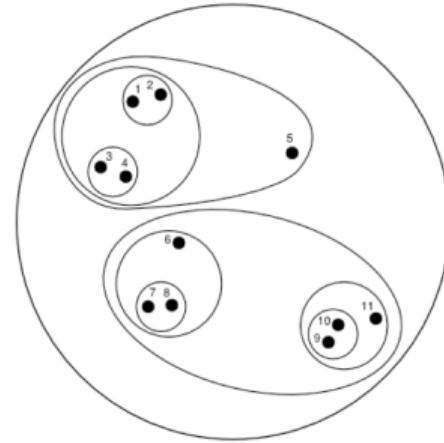
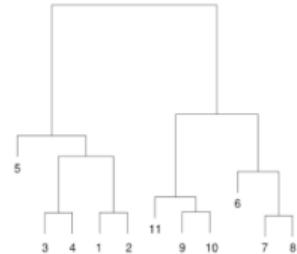
Hierarchical clustering

Generates a **nested structure** of clusters

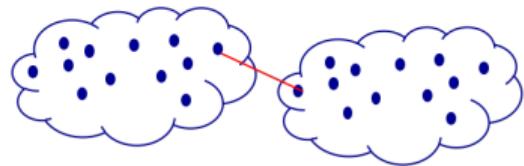
- Agglomerative (bottom up)
 - as a starting state, each data point is a cluster
 - in each step the two **less separated** clusters are merged into one
 - a measure of **separation between clusters** is needed
- Divisive (top down)
 - as a starting state, the entire dataset is the only cluster
 - in each step, the cluster with the lowest cohesion is split
 - a measure of cluster cohesion and a split procedure are needed

Hierarchical clustering output

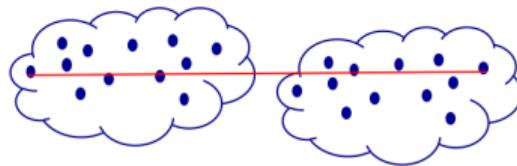
- Dendrogram (left)
- Nested cluster diagram (right)
- They represent the same structure
- The representation is the same for agglomerative and divisive
- The agglomerative methods are the most used



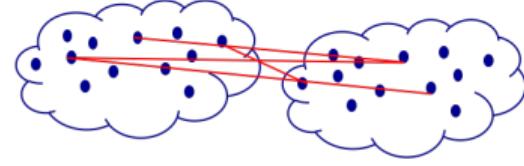
Separation between clusters – Graph based



Single Link



Complete Link



Average Link

$$\text{Sep}(k_i, k_j) = \min_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

$$\text{Sep}(k_i, k_j) = \max_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

$$\text{Sep}(k_i, k_j) = \frac{1}{|k_i||k_j|} \sum_{x \in k_i, y \in k_j} \text{Dist}(x, y)$$

The distance between sets is based on the distances between objects belonging to the two sets, respectively

Separation between clusters – Ward's method

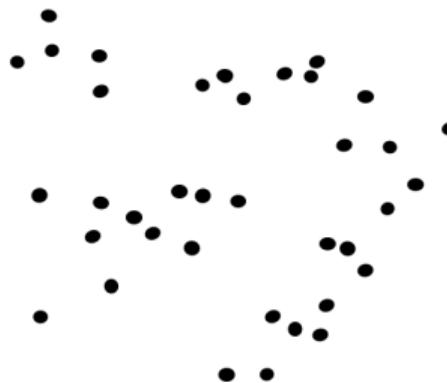
Given two sets S_1 and S_2 with the respective $SSE(S_1)$ and $SSE(S_2)$, the distance is computed as

$$d(S_1, S_2) = SSE(S_1 \cup S_2) - (SSE(S_1) + SSE(S_2))$$

Smaller separation implies a lower increase in the SSE after merging

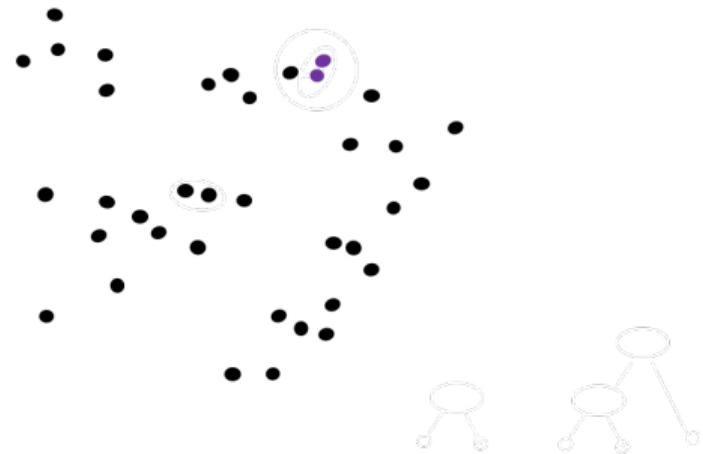
Single linkage hierarchical clustering I

1. Initialization: every object is a cluster



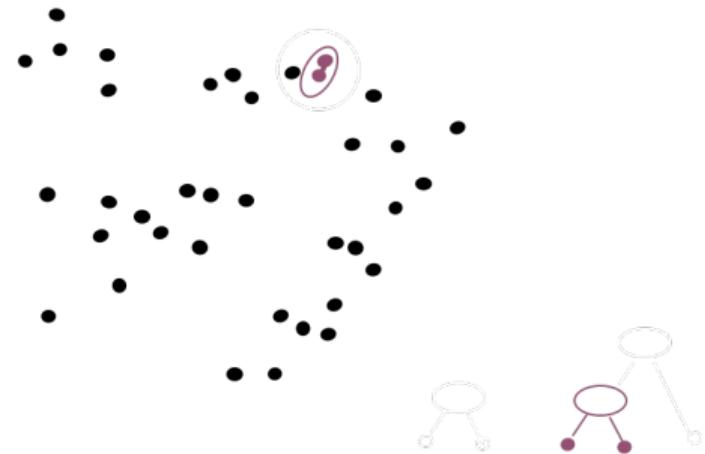
Single linkage hierarchical clustering II

1. Initialization: every object is a cluster
2. Find the **less separated** pair



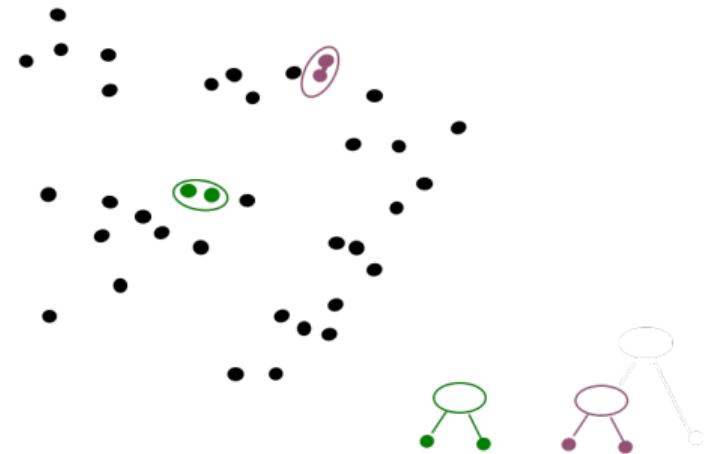
Single linkage hierarchical clustering III

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster



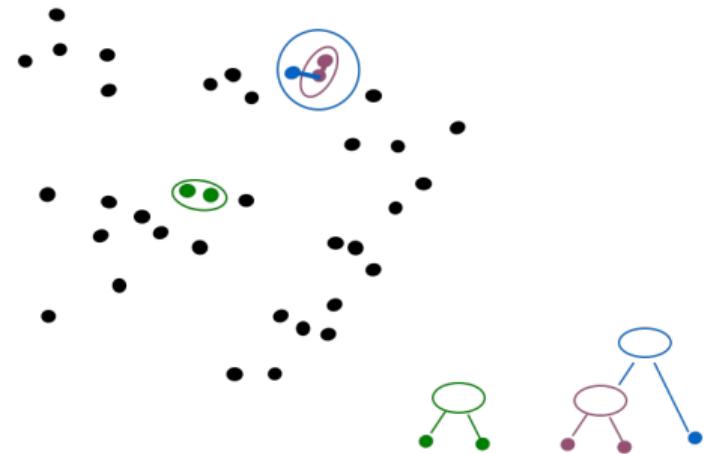
Single linkage hierarchical clustering IV

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat



Single linkage hierarchical clustering V

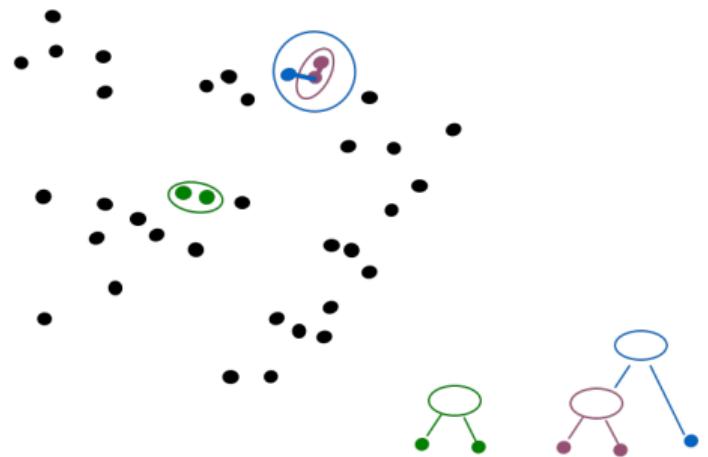
1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat ...



Single linkage hierarchical clustering VI

1. Initialization: every object is a cluster
2. Find the **less separated** pair
3. Merge in a single cluster
4. Repeat ...

The result is a **dendrogram** (taxonomy, object hierarchy)



Single linkage algorithm

- Initialize the clusters, one for each objects
- Compute the **distance matrix** between the clusters, squared, symmetric, the size is the number of objects N , the main diagonal is null
- While the number of clusters is greater than 1
 - find the two clusters with lowest separation, say k_r and k_s
 - merge them in a cluster
 - delete from the distance matrix the rows and columns r and s and insert one new row and column with the distances of the new cluster from the others

$$\text{Dist}(k_k, k_{(r+s)}) = \min(\text{Dist}(k_k, k_r), \text{Dist}(k_k, k_s)) \forall k \in [1, K]$$

Time and space complexity

- Space and time: $\mathcal{O}(N^2)$ for the computation and the storage of the distance matrix
- Worst case $N - 1$ iterations to reach the final single cluster
- For the i -th step of the main iteration:
 - search of the pair to merge $\mathcal{O}((N - i)^2)$
 - recomputation of the distance matrix $\mathcal{O}((N - i))$
- Time, in summary: $\mathcal{O}(N^3)$
 - can be reduced to $\mathcal{O}(N^2 \log(N))$ with indexing structures

Italian cities example I

	BA	FI	MI	NA	RM	TO
BA	0	662	877	255	412	996
FI	662	0	295	468	268	400
MI	877	295	0	754	564	138
NA	255	468	754	0	219	869
RM	412	268	564	219	0	669
TO	996	400	138	869	669	0



Italian cities example II

	BA	FI	MI	NA	RM	TO
BA	0	662	877	255	412	996
FI	662	0	295	468	268	400
MI	877	295	0	754	564	138
NA	255	468	754	0	219	869
RM	412	268	564	219	0	669
TO	996	400	138	869	669	0



Italian cities example III

	BA	FI	MI/TO	NA	RM
BA	0	662	877	255	412
FI	662	0	295	468	268
MI/TO	877	295	0	754	564
NA	255	468	754	0	219
RM	412	268	564	219	0



Italian cities example IV

	BA	FI	MI/TO	NA/RM
BA	0	662	877	255
FI	662	0	295	268
MI/TO	877	295	0	564
NA/RM	255	268	564	0



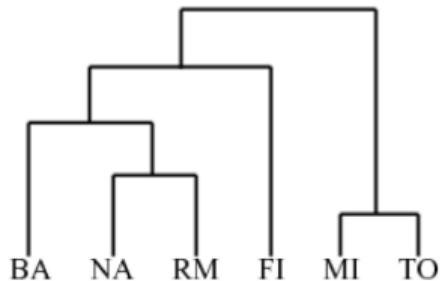
Italian cities example V

	BA/NA/RM	FI	MI/TO
BA/NA/RM	0	268	564
FI	268	0	295
MI/TO	564	295	0

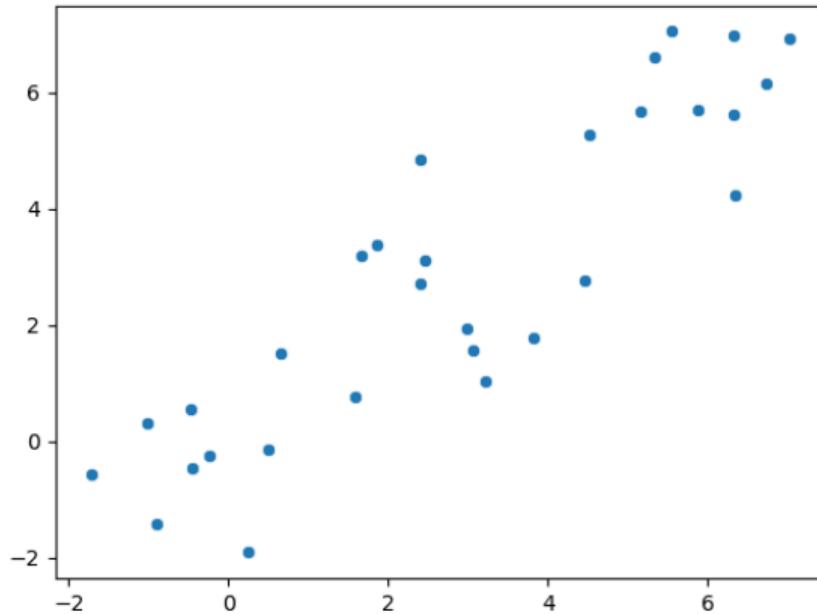


Italian cities example VI

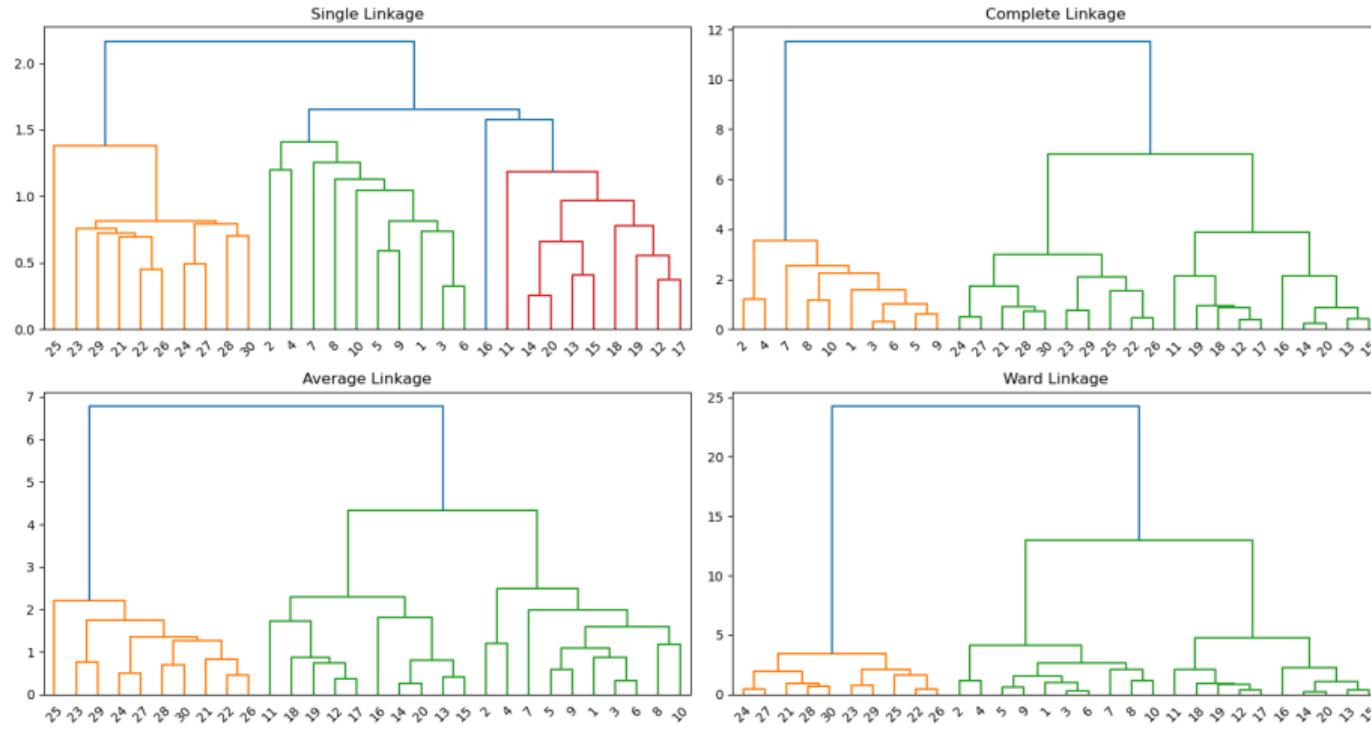
	BA/FI/NA/RM	MI/TO
BA/FI/NA/RM	0	295
MI/TO	295	0



Comparison of linking methods - Sample data

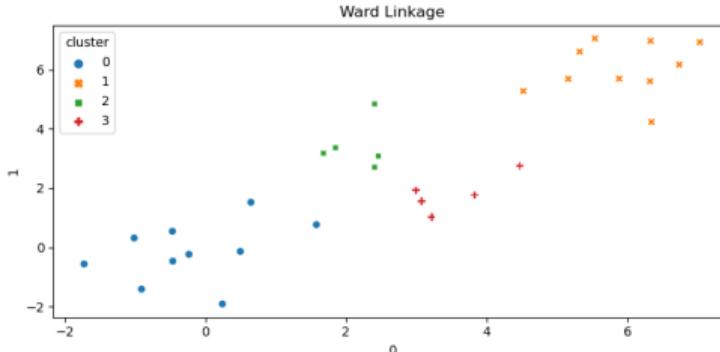
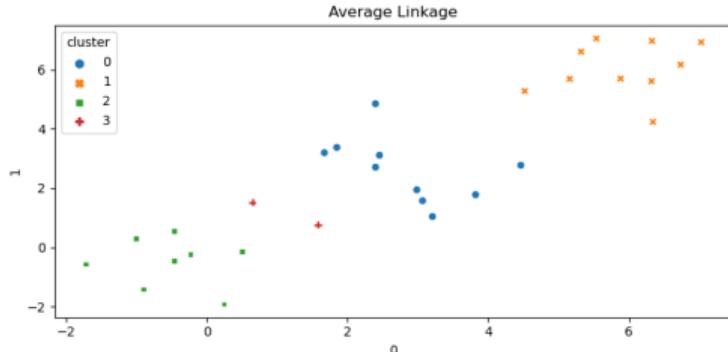
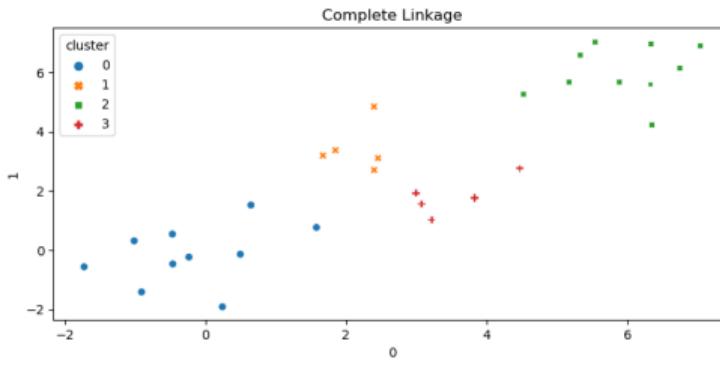
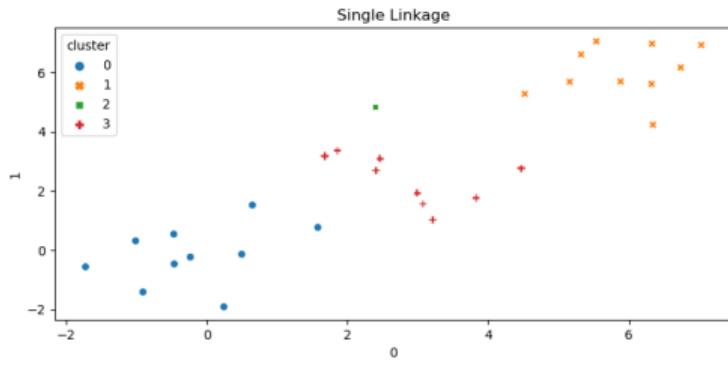


Comparison of linking methods - Dendograms



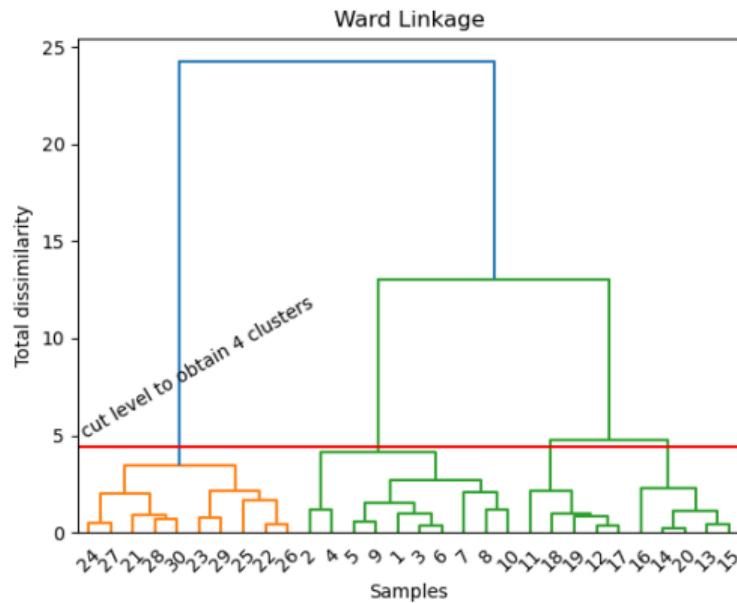
Comparison of linking methods

Labelled data with 4 clusters



Generating the clustering scheme

- The desired clustering scheme is obtained by choosing the number of clusters
- The choice of the level is application dependent, and can also be guided by indexes, as in the case of K-means
- The choice is equivalent to **cutting** the dendrogram at the appropriate level



Discussion I

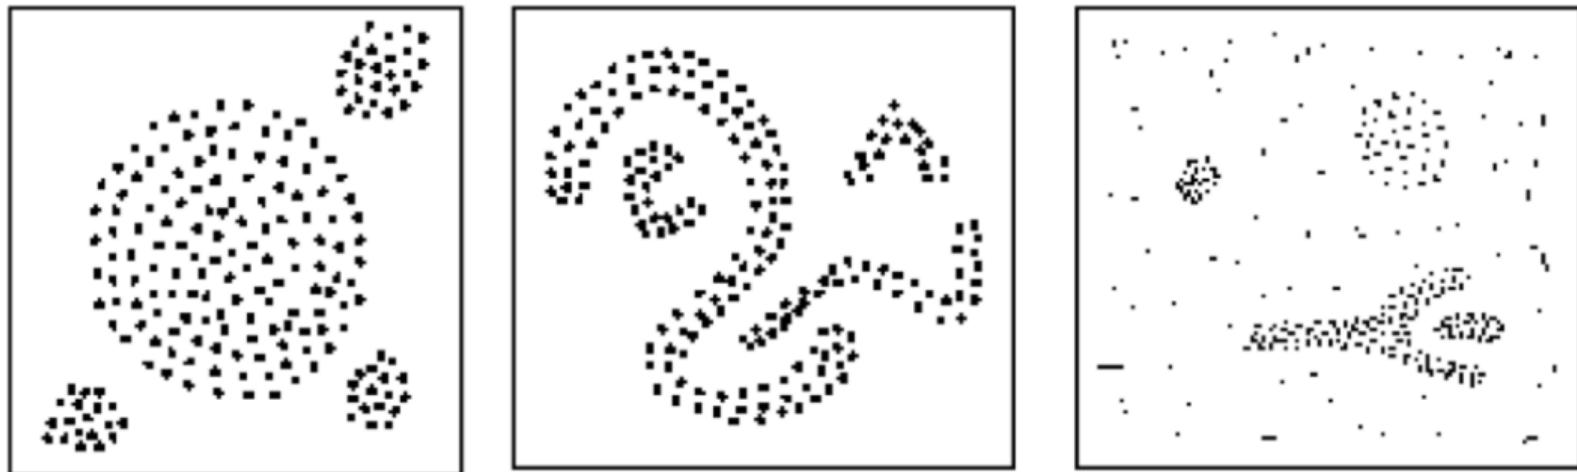
- The vertical axis in the dendrogram is the **total dissimilarity** inside the clusters, which obviously increases for decreasing number of clusters
- The **diameter** of a cluster is the distance among the most separated objects
 - Single linkage tends to generate clusters with larger diameters also at low levels
 - Complete linkage tends to generate more compact clusters
 - Ward linkage tends to generate more *spheric* clusters

Discussion II

- (?): The scaling is poor, due to the high complexity
- (?): There isn't a global objective function, the decision is always local and cannot be undone
- (?): The dendrogram structure is of great help for the interpretation of the result
- (?): Empirically, the result is frequently good

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
5	Density based clustering	94
	● DBSCAN	97
	● Kernel Density Estimation	110
6	Model based clustering	114
7	Final remarks	123

Density based clustering



Clusters are high-density regions separated by low-density regions

Computing density

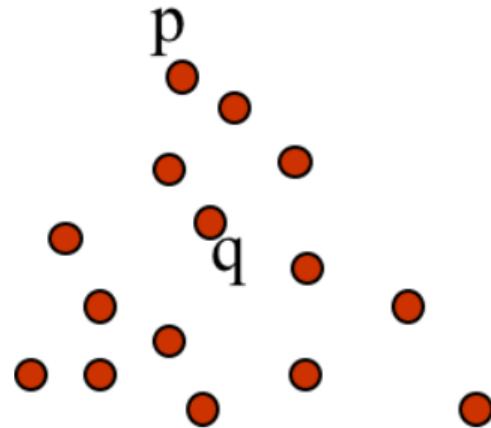
The two most obvious solutions

- Grid-based
 - split the (hyper)space into a regularly spaced grid
 - count the number of objects inside each grid element
- Object-centered
 - define the radius of a (hyper)sphere
 - attach to each object the number of objects which are inside that sphere

DBSCAN – Density Based Spatial Clustering of Applications with Noise²

Intuition

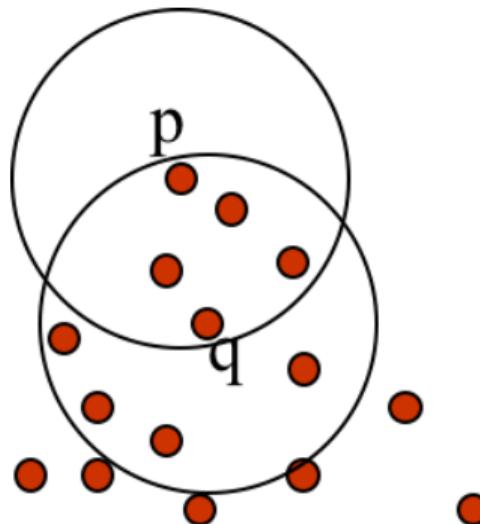
- intuitively, p is a **border** point, while q is a **core** point



2 [Ester et al.(1996)] Ester, Kriegel, Sander, and Xu]

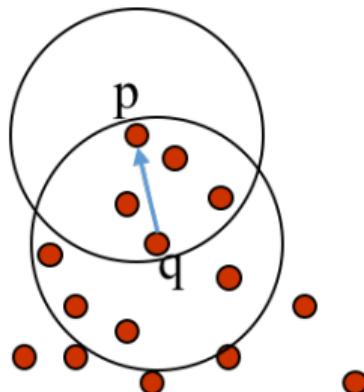
Neighborhood

- define a radius ϵ and define as **neighborhood** of a point the ϵ -hypersphere centered at that point
- points p and q are one in the neighborhood of the other
 - neighborhood is *symmetric*



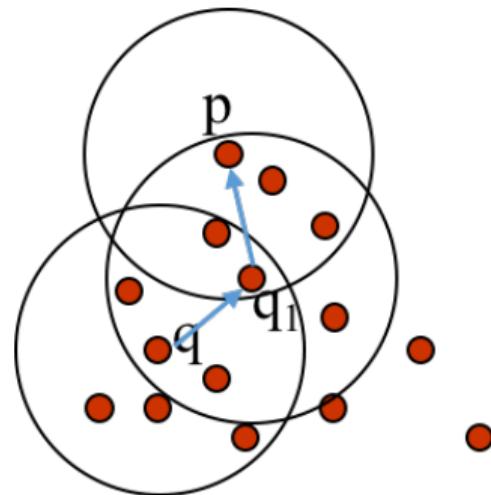
Direct Density Reachability

- define a threshold `minPoints` and define as **core** a point with at least `minPoints` points in its neighborhood, as **border** otherwise
 - with $\text{minPoints} = 5$, q is core, p is border
- define that a point p is **directly density reachable** from point q iff
 - q is core
 - q is in the neighborhood of p
- direct density reachability is not symmetric
 - in the example q is not directly density reachable from p , since p is border



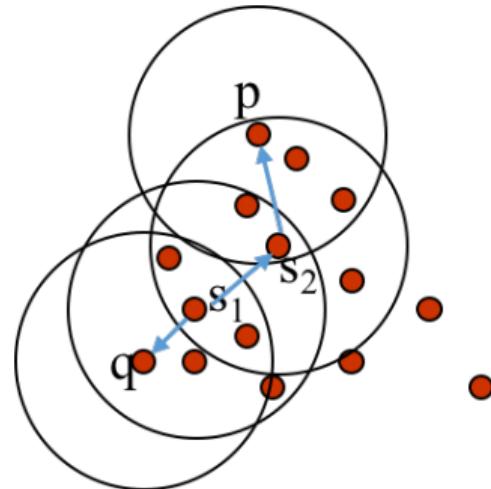
Density Reachability

- a point p is **density reachable** from point q iff
 - q is core
 - there is a sequence of points point q_i such that q_{i+1} is directly density reachable from q_i , $i \in [1, nq]$, q_1 is directly reachable from q and p is directly density reachable from q_{nq}
- reachability is not symmetric
 - in the example q is not density reachable from p , since p is border



Density Connection

- a point p is **density connected** to point q iff there is a point s such that p and q are density reachable from s
- density connection is symmetric



Generation of clusters

- A *cluster* is a maximal set of points connected by *density*
- Border points which are not connected by density to any core point are labelled as **noise**

Algorithm I

Algorithm 1 DBSCAN

Require: SetOfPoints: UNCLASSIFIED points

Require: Eps, MinPts

ClusterId \leftarrow nextId(NOISE);

for i = 1 to SetOfPoints.size **do**

 Point \leftarrow SetOfPoints.get(i)

if Point.CId = UNCLASSIFIED **then**

if ExpandCluster(SetOfPoints, Point, ClusterId, Eps, MinPts)

then

 ClusterId \leftarrow nextId(ClusterId)

Ensure: SetOfPoints

Algorithm II

```
ExpandCluster(SetOfPoints, Point, ClId, Eps, MinPts) : Boolean;  
    seeds:=SetOfPoints.regionQuery(Point,Eps);  
    If seeds.size<MinPts THEN // no core point  
        SetOfPoint.changeClId(Point,NOISE);  
        RETURN False;  
    ELSE ...
```

Algorithm III

```
// all points in seeds are density-reachable from Point
SetOfPoints.changeC1Ids(seeds,C1Id);
seeds.delete(Point);
WHILE seeds <> Empty DO
    currentP := seeds.first();
    result := SetOfPoints.regionQuery(currentP,Eps);
    If result.size >= MinPts THEN
        For i FROM 1 TO result.size DO
            resultP := result.get(i);
            If resultP.C1Id IN {UNCLASSIFIED, NOISE} THEN
                If resultP.C1Id = UNCLASSIFIED THEN seeds.append(resultP);
            END If;
            SetOfPoints.changeC1Id(resultP,C1Id);
        END If; // UNCLASSIFIED or NOISE
        END For;
    END If; // result.size >= MinPts
    seeds.delete(currentP);
END WHILE; // seeds <> Empty
RETURN True;
END If
END; // ExpandCluster
```

How to set ϵ and minPoints?

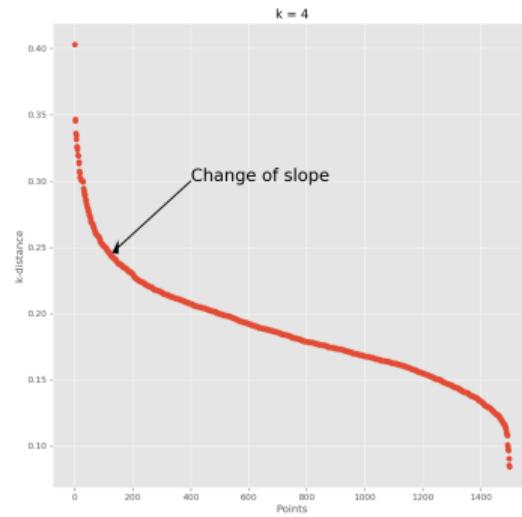
- As in many other machine learning algorithms, a *grid search* over several combination of hyperparameters can be useful
- As a *rule of thumb*, you can try $\text{minPoints} = 2 * D$, the number of dimensions
- Noise suggest an increase in minPoints
- A guess for ϵ requires more effort, considering the distance of the k -nearest neighbour, with $k = \text{minPoints}$

Good guess for ϵ !

- Consider the vector of the *k-distances*
 - choose k
 - for each point we compute the distance of its k -nearest neighbour and we sort the points for decreasing k -distance
- Choosing a given k -distance as ϵ , it turns out that all the points with a k -distance bigger than ϵ will be considered as *border*
 - in the figure of next page they are the points to the left of the vertical of the chosen ϵ

Good guess for $\epsilon \parallel$

- Usually, datasets which exhibit some tendency to clustering exhibit also a *change of slope*
- The best ϵ can be found with a grid search in the *area of the change of slope*
 - the figure refers to a dataset with 1500 points and with minPoints=4
 - this figure suggests a fine tuning of ϵ in the interval 0.2–0.3



Comments

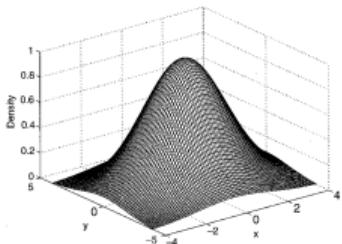
- ☺ Finds clusters of any shape
- ☺ Is robust w.r.t. noise
- ☹ Problems if clusters have widely varying densities
 - Being based on distances between points, the complexity is $\mathcal{O}(N^2)$
 - reduced to $\mathcal{O}(N \log(N))$ if spatial indexes, such as R*, are available
 - Very sensitive to the values of ϵ and minPoints
 - Decreasing ϵ and increasing minPoints reduces the cluster size and increases the number of noise points



Kernel Density Estimation 3

- A technique developed in statistics and pattern mining
- Describe the distribution of the data by a function
- The overall density function is the sum of the **influence functions** (or **kernel functions**) associated with each point
- The kernel function
 - must be **symmetric** and monotonically decreasing
 - usually has a *parameter* to set the decreasing rate

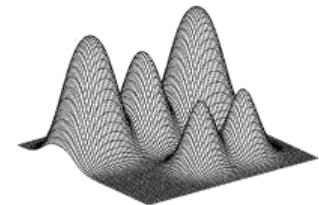
KDE – Example



Gaussian kernel



Overall density – Grey scale plot



Sample set of 12 points
Overall density – Surface plot

DENCLUE algorithm

1. Derive a density function for the space occupied by the data points
2. Identify the points that are local maxima
3. Associate each point with a density attractor by moving in the directions of maximum increase in density
4. Define clusters consisting of points associated with a particular density attractor
5. Discard clusters whose density attractor has a density less than a user-specified threshold ξ
6. Combine clusters that are connected by a path of points that all have a density of ξ or higher

DENCLUE comments

- ☺ It has a strong theoretical foundation on statistics
 - precise computation of density
 - DBSCAN is a special case of DENCLUE where the influence is a step function
- ☺ Good at dealing with noise and clusters of different shapes and sizes
- ☹ expensive computation $\mathcal{O}(N^2)$
 - can be optimized with approximated *grid based* computation
- ☹ Troubles with high dimensional data and clusters with different densities

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
5	Density based clustering	94
6	Model based clustering	114
	● Gaussian Mixture	116
7	Final remarks	123

Model based (or statistic based) clustering⁴

- Estimate the parameters of a statistical model to maximize the ability of the model to **explain the data**
- The main technique is to use the **mixture models**
 - view the data as a set of observation from a mixture of different probability distributions
- Usually, the base model is a multivariate normal
 - well-known, easy to work with, good results
- The estimation is usually done using the **maximum likelihood**
 - given a set of data \mathcal{X} , the probability of the data, regarded as a function of the parameters, is called a **likelihood function**
- Attributes are assumed to be random independent variables

4 [Tan et al.(2006) Tan, Steinbach, and Kumar], Section 9.2.2

Gaussian Mixture

a.k.a. Expectation Maximization – EM

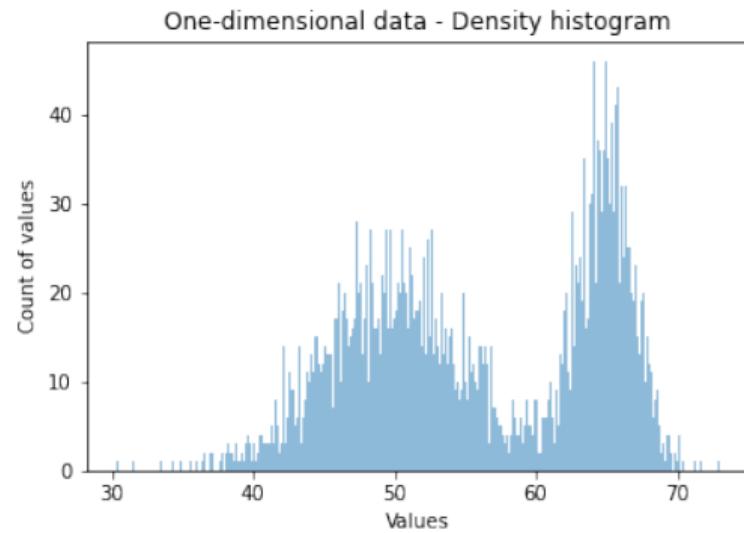
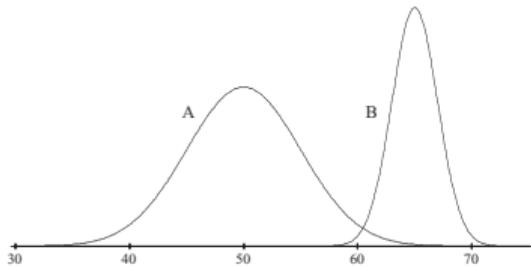
- If the data can be approximated by a single distribution, the derivation of the parameters is straightforward
- In the general case, with many mixed distributions, the EM algorithm is used

EM algorithm

1. Select an initial set of model parameters
2. **repeat**
 - 2.1 **Expectation Step** – For each object, calculate the probability that each object belongs to each distribution
 - 2.2 **Maximization Step** – Given the probabilities from the expectation step, find the new estimates of the parameters that maximize the expected likelihood
3. **until** – the parameters do not change (or the change is below a specified threshold)

One dimension mixture example

- Case with one dimension, two components
- Synthetic data randomly generated with two gaussians
 $\mu_A = 50, \sigma_A = 5, p_A = 0.6$
 $\mu_B = 65, \sigma_B = 2, p_B = 0.4$



EM – one dimension, two clusters example I

- Need to estimate 5 parameters
 - mean and standard deviation for cluster A
 - mean and standard deviation for cluster B
 - sampling probability p for cluster A

$$\Pr(A|x) = \frac{\Pr(x|A)\Pr(A)}{\Pr(x)} = \frac{f(x; \mu_A, \sigma_A)p_A}{\Pr(x)}$$
$$f(x; \mu_A, \sigma_A) = \frac{1}{\sqrt{2\pi}\sigma} e^{\frac{(x-\mu)^2}{2\sigma^2}}$$

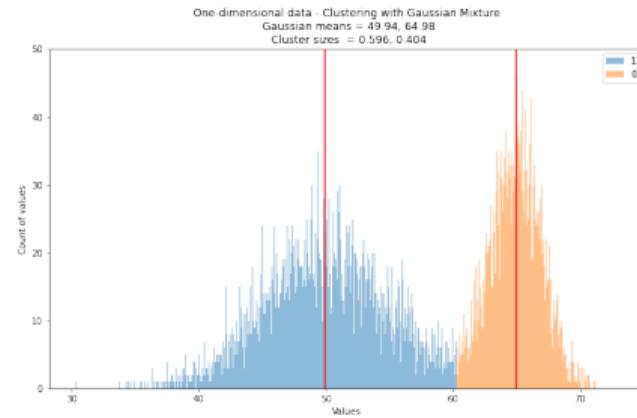
EM – one dimension, two clusters example II

- Repeat until convergence
 - Expectation: Compute p_A and p_B using the current distribution parameters
 - Compute the numerators for $\Pr(A|x)$ and $\Pr(B|x)$ and normalize dividing by their sum
 - Maximization of the distribution likelihood given the data
 - Compute the new distributions parameters, weighting the probabilities according to the current distribution parameters
- After convergence label each object with A or B according to the maximum probability, given the last distribution parameters

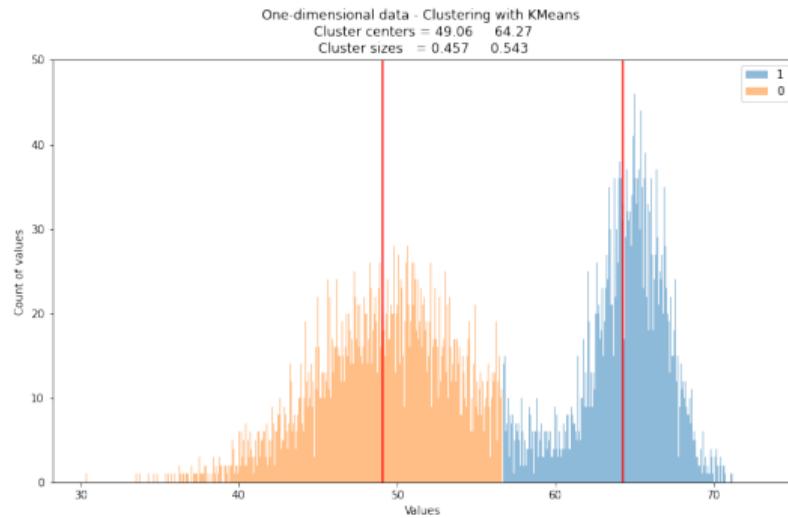
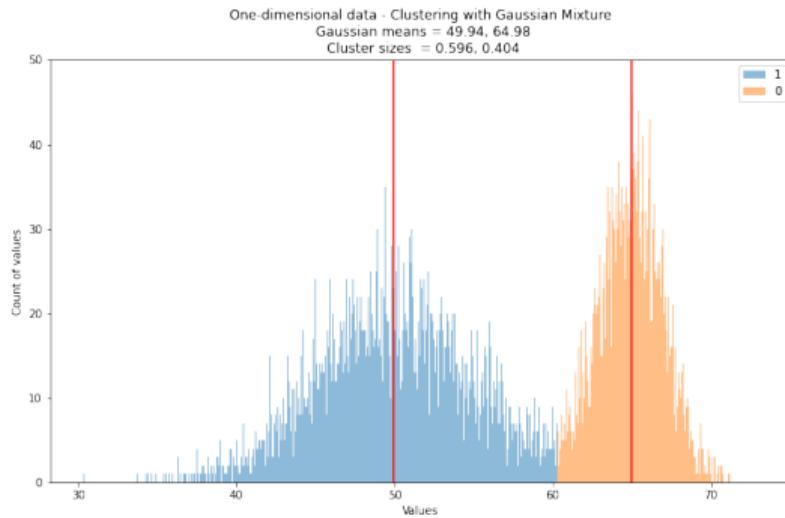
One dimension example - Gaussian Mixture result

Estimated parameters of the distributions

	weight	mean	deviation
0	0.4037	64.978754	2.030968
1	0.5963	49.939990	4.999235



Gaussian Mixture and KMeans – Comparison of results

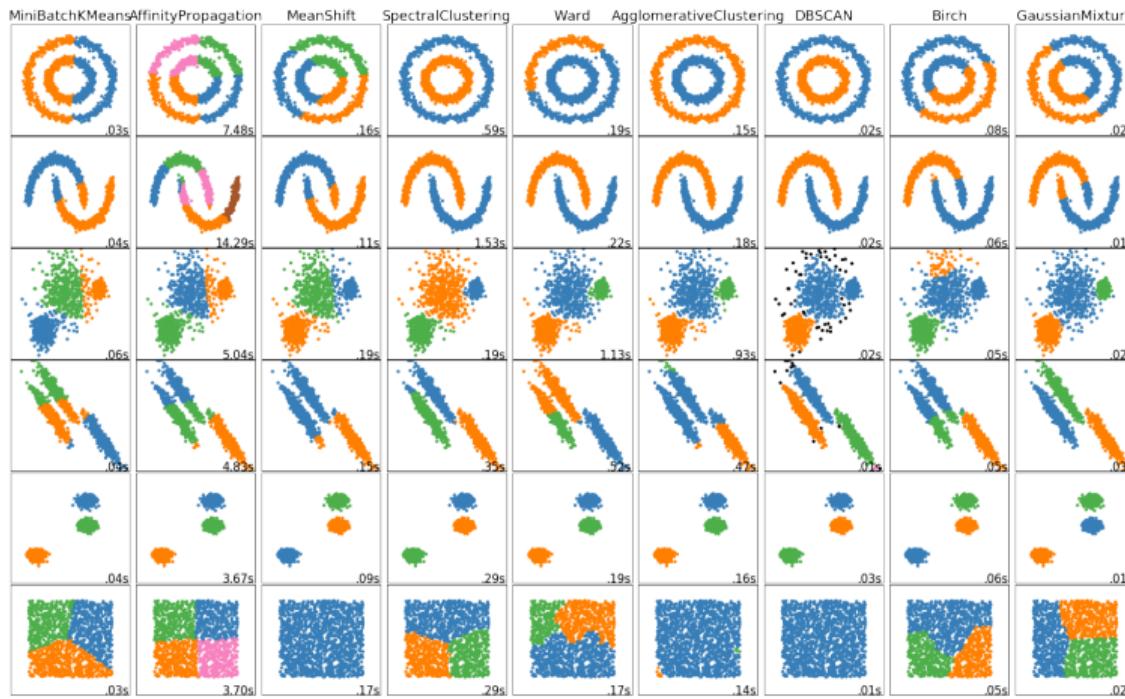


- These data have a *bimodal, gaussian-like* distribution
- The EM algorithm is founded on the hypothesis of modelling data with gaussians
- KMeans is *non-parametric*, and in this case the performance is *worse*

1	Introduction to clustering	2
2	K-means	12
3	Evaluation of a clustering scheme	47
4	Hierarchical clustering	69
5	Density based clustering	94
6	Model based clustering	114
7	Final remarks	123

Comparison of results for selected algorithms

(from [Scikit-Learn documentation](#))



A summary of selected clustering algorithms – I

Method name	Parameters	Scalability	Use case	Geometry (metric used)
K-Means	number of clusters	Very large n_samples, medium n_clusters with MiniBatch code	General-purpose, even cluster size, flat geometry, not too many clusters, inductive	Distances between points
Affinity propagation	damping, sample preference	Not scalable with n_samples	Many clusters, uneven cluster size, non-flat geometry, inductive	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	bandwidth	Not scalable with n_samples	Many clusters, uneven cluster size, non-flat geometry, inductive	Distances between points
Spectral clustering	number of clusters	Medium n_samples, small n_clusters	Few clusters, even cluster size, non-flat geometry, transductive	Graph distance (e.g. nearest-neighbor graph)
Ward hierarchical clustering	number of clusters or distance threshold	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints, transductive	Distances between points

A summary of selected clustering algorithms – II

Method name	Parameters	Scalability	Use case	Geometry (metric used)
Agglomerative clustering	number of clusters or distance threshold, linkage type, distance	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints, non Euclidean distances, transductive	Any pairwise distance
DBSCAN	neighborhood size	Very large n_samples, medium n_clusters	Non-flat geometry, uneven cluster sizes, outlier removal, transductive	Distances between nearest points
OPTICS	minimum cluster membership	Very large n_samples, large n_clusters	Non-flat geometry, uneven cluster sizes, variable cluster density, outlier removal, transductive	Distances between points
Gaussian mixtures	many	Not scalable	Flat geometry, good for density estimation, inductive	Mahalanobis distances to centers
BIRCH	branching factor, threshold, optional global clusterer.	Large n_clusters and n_samples	Large dataset, outlier removal, data reduction, inductive	Euclidean distance between points

Clustering types

- Partitioning
 - iteratively find partitions in the dataset, optimizing some quality criterion
- Hierarchic
 - recursively compute a structured hierarchy of subsets
- Density based
 - compute densities and aggregates clusters in high density areas
- Model based
 - assume a model for the distribution of the data and find the model parameters which guarantee the best fitting to the data

Clustering scalability

- Effectiveness decreases with
 - dimensionality D
 - noise level
- Computational cost increases with
 - dataset size N , at least linearly
 - dimensionality D

Research perspective

- From the past
 - well-known problem in statistics
 - recent research
 - machine learning
 - databases
 - visualization
- For the future
 - effective and efficient algorithms for big data clustering, with a large number of dimensions, noise requested scalability w.r.t.:
 - size (N)
 - dimensionality (D)
 - noise level
 - rate of new data arrivals

Uses of clustering – data comprehension

- Biology
 - Creation of taxonomies
 - Genetics
- Information Retrieval
 - Grouping documents
- Climatology
 - Repetition patterns
- Psychology and medicine
 - Identification of illness types in front of partial variation of evidence
- Business
 - Customer grouping

Uses of clustering – utilities

- Summarization
 - Reasoning with groups representatives instead of with the entire population
- Data compression
 - Reduce the amount of data
 - Find cluster prototypes and substitute data with the indexes of the prototypes
 - Vector quantization
- Find the nearest neighbours
 - Each object refers to his prototypes
 - Near neighbours refer to the same prototype

Bibliography I

- Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu.
A density-based algorithm for discovering clusters in large spatial databases with noise.
pages 226–231. AAAI Press, 1996.
- A. K. Jain, M. N. Murty, and P. J. Flynn.
Data clustering: A review.
ACM Comput. Surv., 31(3):264–323, September 1999.
ISSN 0360-0300.
doi: 10.1145/331499.331504.
URL <http://doi.acm.org/10.1145/331499.331504>.

Bibliography II

- ▶ Pang-Nin Tan, Michael Steinbach, and Vipin Kumar.
Data Mining.
Addison Wesley, 2006.
ISBN 0-321-32136-7.